

# TensorFlow Tutorial

## Strata NYC 2016

# Your guides



Martin Wicke  
[@martin\\_wicke](https://twitter.com/martin_wicke)



Josh Gordon  
[@random\\_forests](https://twitter.com/random_forests)

# Slides and code

[goo.gl/nrdsxM](http://goo.gl/nrdsxM)

# Welcome and Logistics

- We will present the tutorial in Jupyter notebooks.
- To run them on your machine, you will need a working TensorFlow installation (v0.10+)
- Follow the instructions at [goo.gl/nrdsxM](http://goo.gl/nrdsxM) to install TensorFlow and Jupyter on your laptop

# Outline

## Part 1) Concepts

- Tensors, graphs, sessions
- Inference, loss, training

## Exercises

- Fibonacci sequence
- Linear regression
- Taxes
- Linear vs Nonlinear Classifiers

## Part 2) Images

- Classifying images
- Image Retraining

## Exercises

- Deep MNIST
- Convolutional MNIST

# Outline cont'd

## Part 3) Text

- Representing words

## Exercises

- Word embeddings
- DBpedia

## Part 4) Free play

- Community contributions

## Experiments

- Deep Dream
- Neural Style



# What's Machine Learning?



# AlphaGo

BBC News Sport Weather iPlayer TV Radio More Search

## Find local news

Home UK World Business Politics Tech Science Health Education Entertainment & Arts More

### Technology

## Google achieves AI 'breakthrough' at Go

An artificial intelligence program developed by Google beats Europe's top player at the ancient Chinese game of Go, about a decade earlier than expected.

© 27 January 2016 | Technology

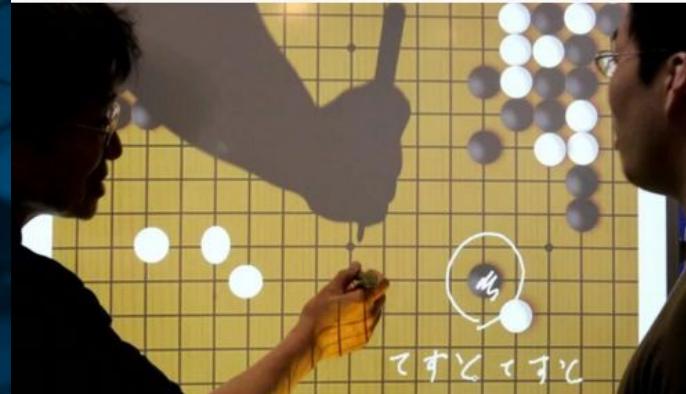
- How did they do it?
- What is the game Go?

Facebook trains AI to beat humans at Go



## Google's AI just cracked the game that supposedly no computer could beat

By Mike Murphy | January 27, 2016



Going up. (Reuters/Kiyoshi Ota)

Computers have slowly started to encroach on activities we previously believed only the brilliantly sophisticated human brain could handle. IBM's Deep Blue supercomputer beat Grand Master Garry Kasparov at chess in 1997, and in 2011 IBM's Watson beat former human winners at the quiz game *Jeopardy*. But the ancient board game Go has long been one of the major goals of artificial intelligence research. It's understood to be one of the most difficult games for computers to handle due to the sheer number of possible moves a player can make at any given point. Until now, that is.



# Open Source Models

[github.com/tensorflow/models](https://github.com/tensorflow/models)

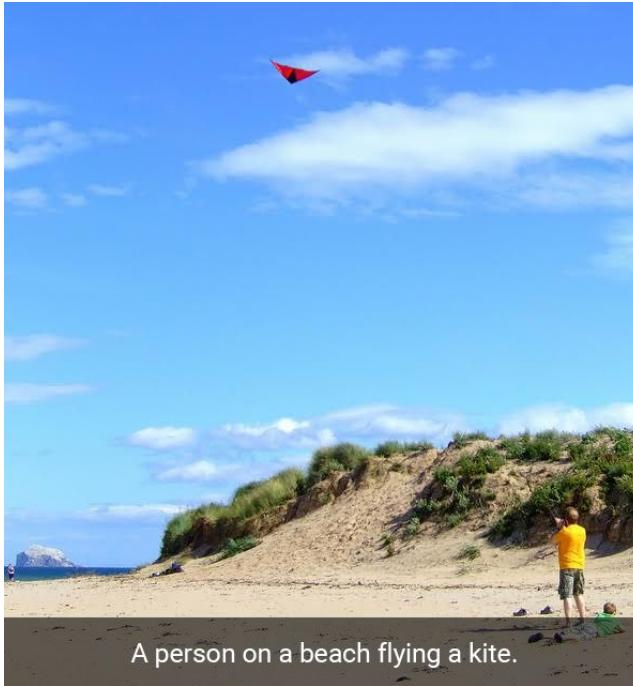
# Inception



An Alaskan Malamute (left) and a Siberian Husky (right). Images from Wikipedia.

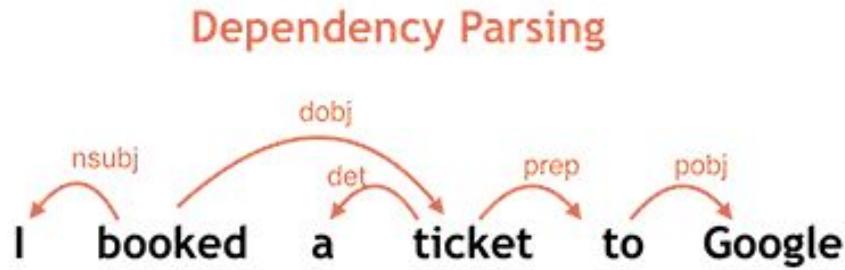
<https://research.googleblog.com/2016/08/improving-inception-and-image.html>

# Show and Tell



<https://research.googleblog.com/2016/09/show-and-tell-image-captioning-open.html>

# Parsey McParseface



<https://research.googleblog.com/2016/05/announcing-syntaxnet-worlds-most.html>

# Text Summarization

## Original text

- *Alice and Bob took the train to visit the zoo. They saw a **baby giraffe**, a **lion**, and a **flock of colorful tropical birds**.*

## Abstractive summary

- *Alice and Bob visited the zoo and saw **animals and birds**.*

# How Can You Get Started?

Three ways, with varying complexity:

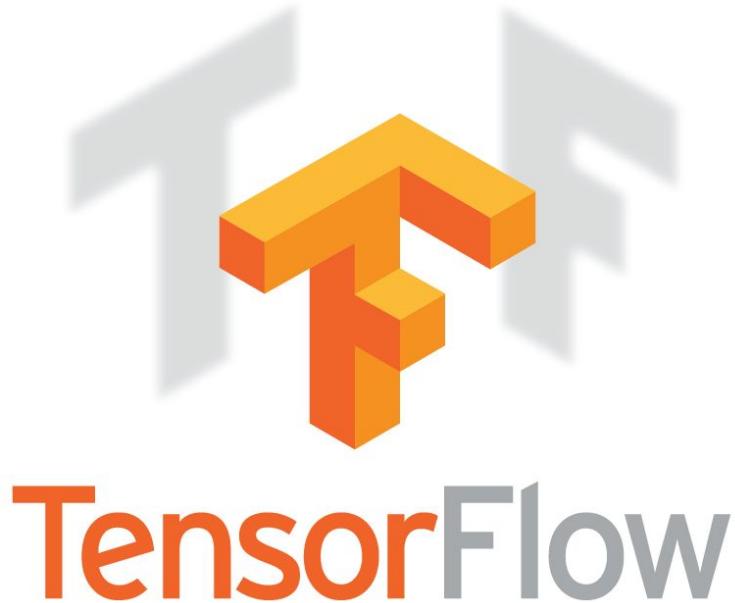
1. Use a Cloud-based API (e.g., Vision)
  
2. Use an existing model architecture, and retrain it or fine tune on your dataset (e.g., Inception)
  
3. Develop your own machine learning models for new problems



More  
flexible,  
but more  
effort



Install TensorFlow  
[goo.gl/nrdsxM](http://goo.gl/nrdsxM)

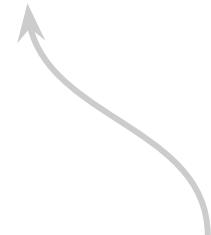


- **Open source** Machine Learning library
- Especially useful for **Deep Learning**
- For research **and** production
- **Apache 2.0** license

A multidimensional array.

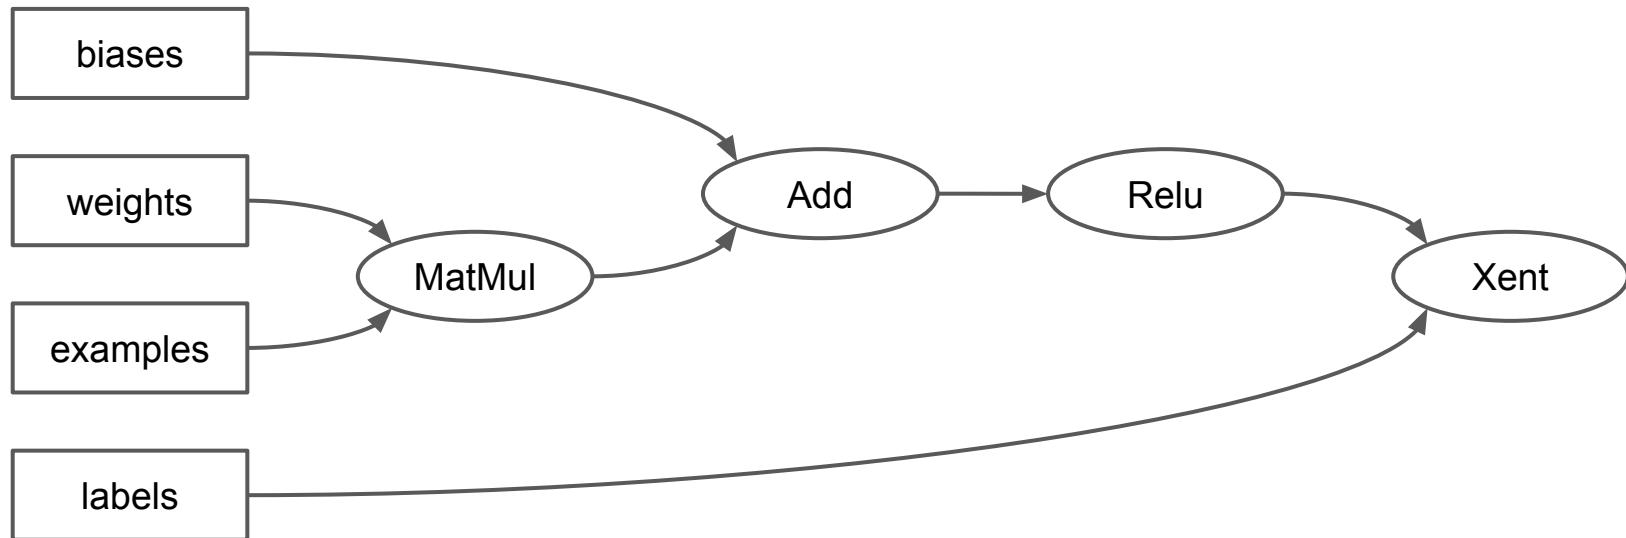


# TensorFlow



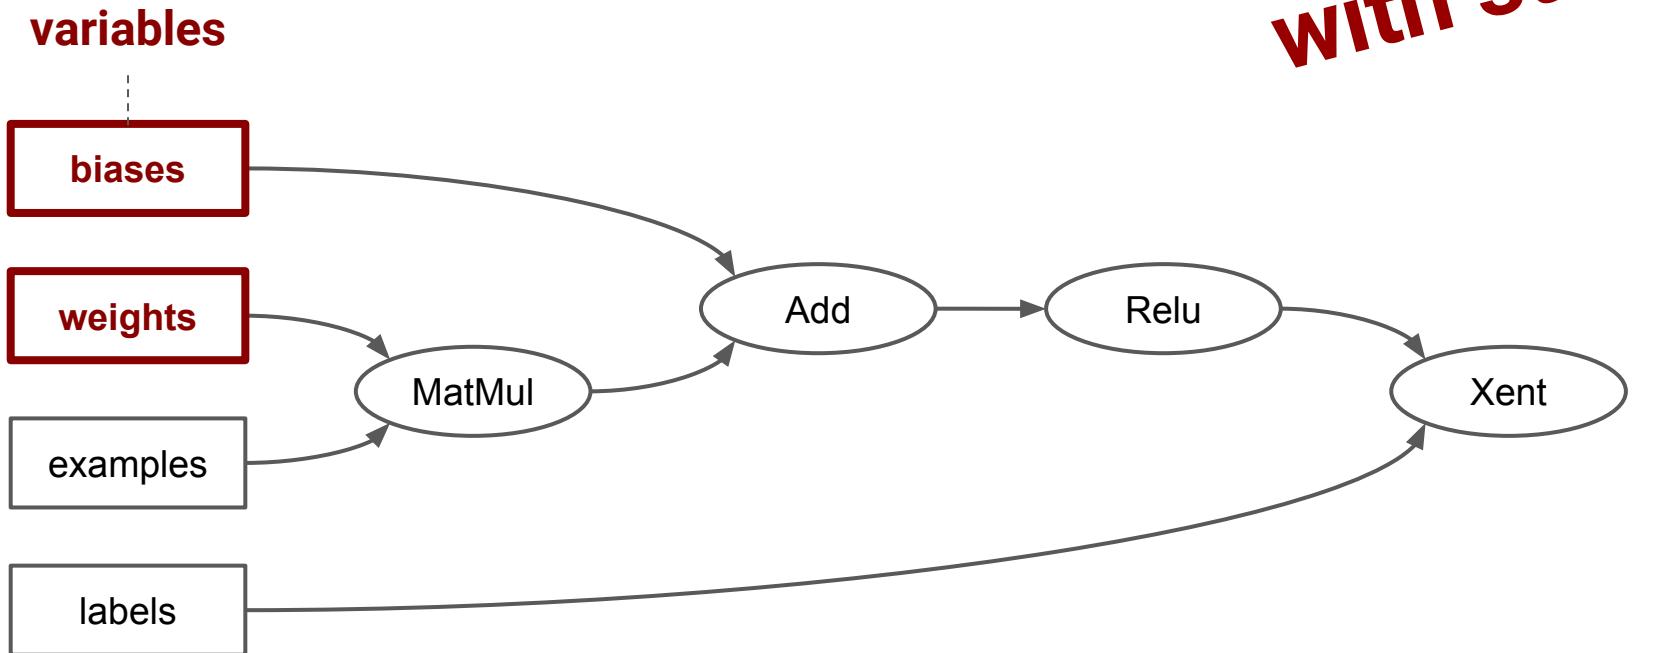
A graph of operations.

# Any Computation is a TensorFlow Graph



# Any Computation is a TensorFlow Graph

*with state*



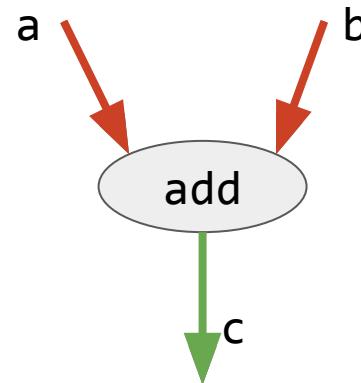
# Build a graph; then run it.

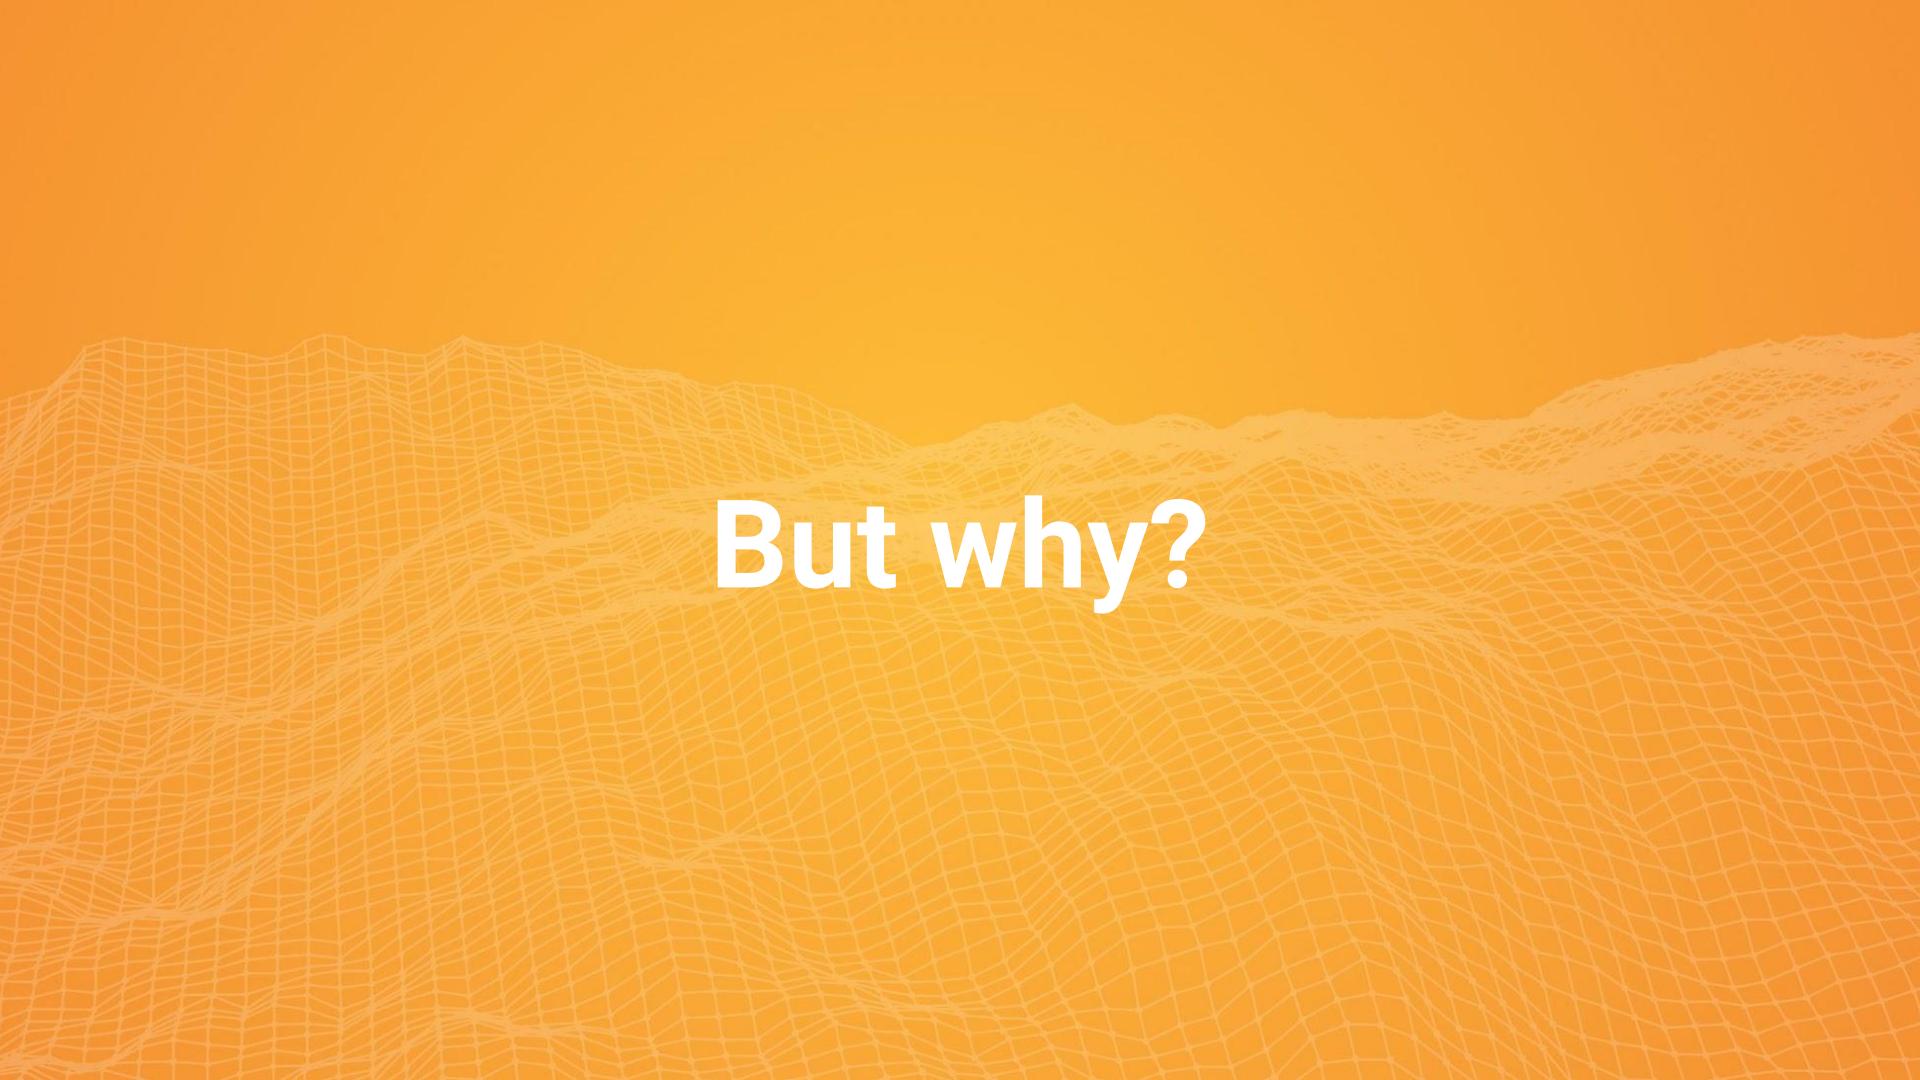
...

```
c = tf.add(a, b)
```

...

```
session = tf.Session()  
value_of_c = session.run(c, {a=1, b=2})
```

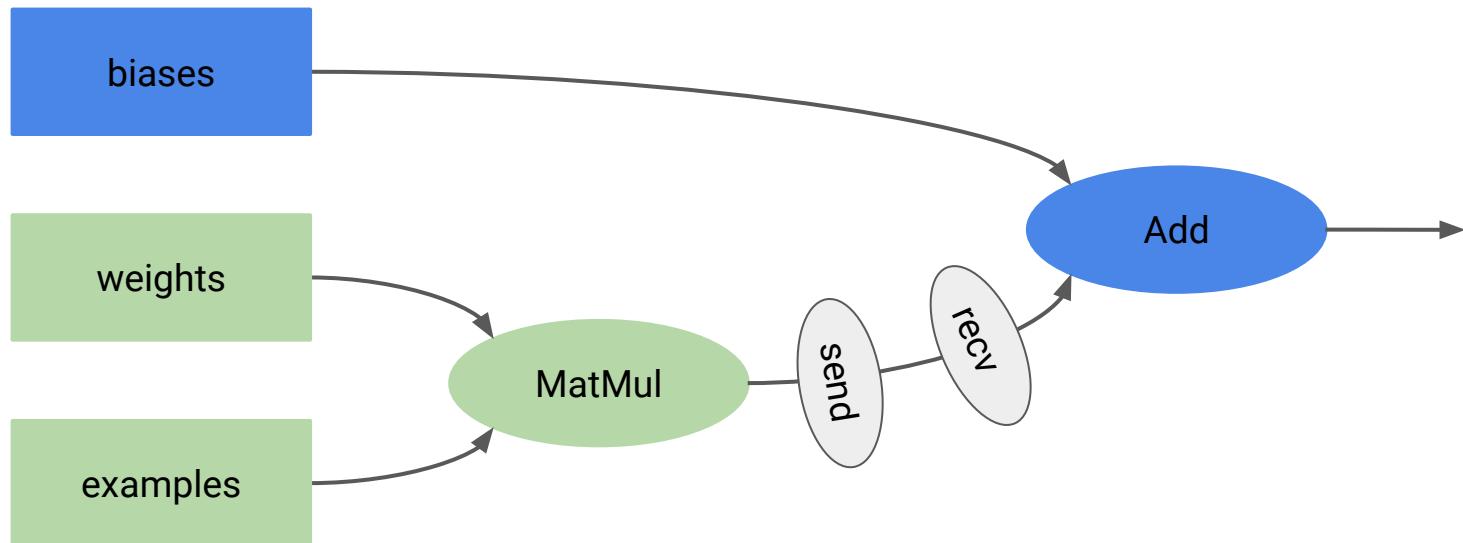


The background features a white wireframe landscape graphic against an orange gradient background. The wireframe consists of a grid of lines forming a series of hills and valleys, with the highest peaks on the left and right sides and lower ridges in the center.

**But why?**

# But, why?

Graphs can be processed, compiled, remotely executed, assigned to devices.



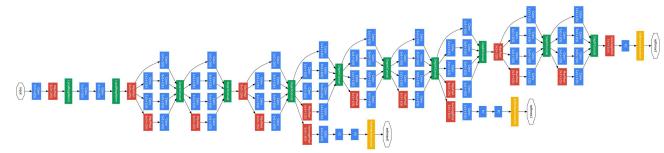
# Three sources of complexity



Heterogenous  
systems

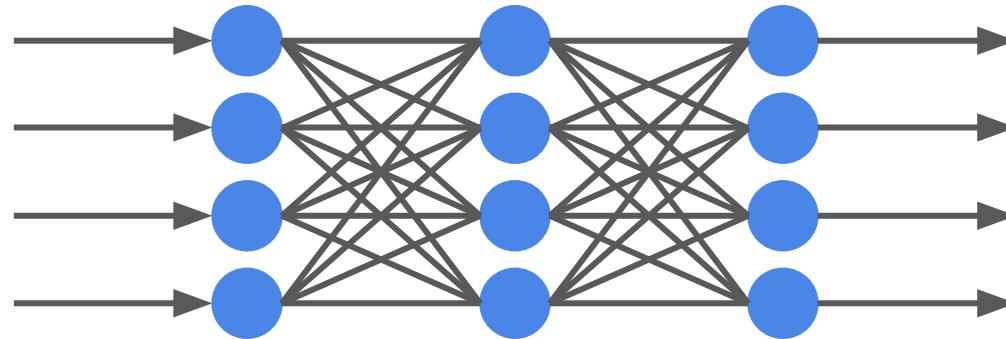


Distributed  
systems

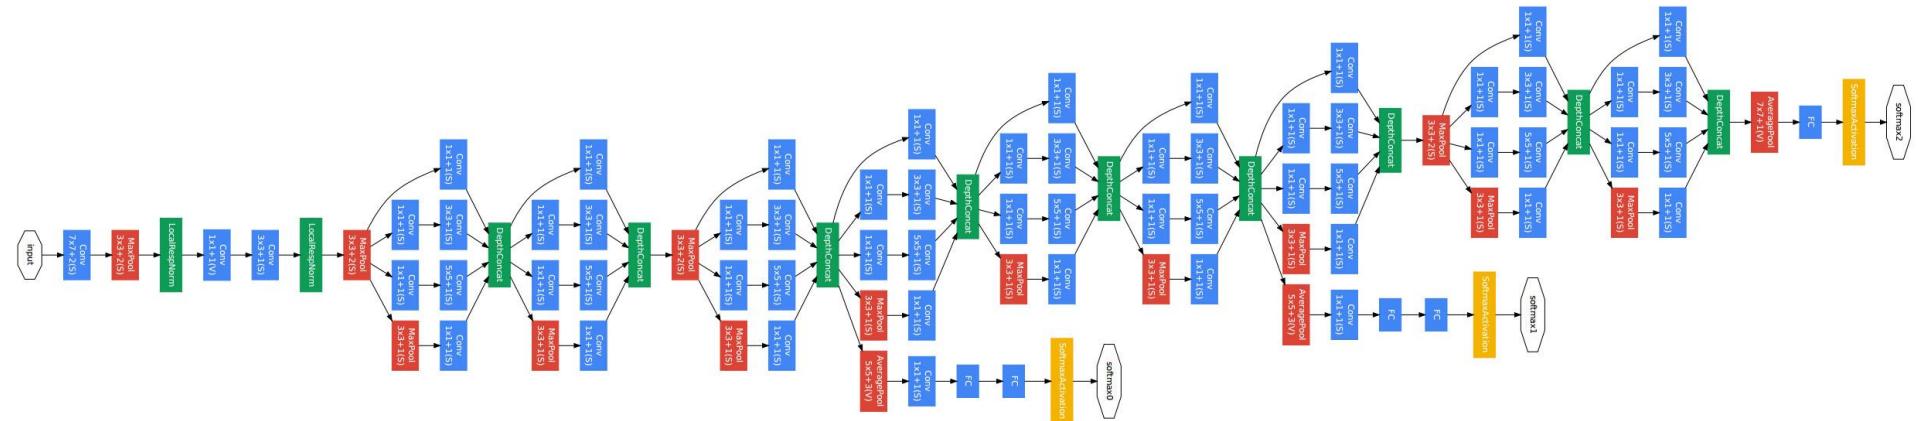


Large models

# Neural networks



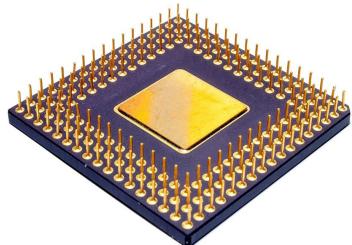
# Inception (2015)



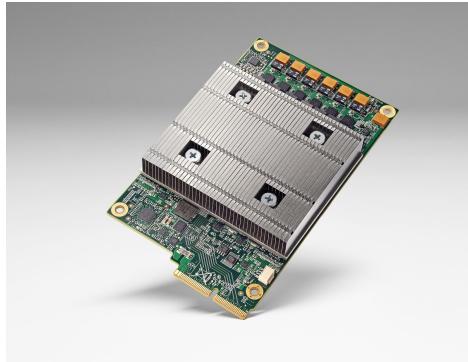
# Distributed systems



# Heterogenous systems



CPU



TPU



GPU



Android



iOS



Raspberry  
Pi

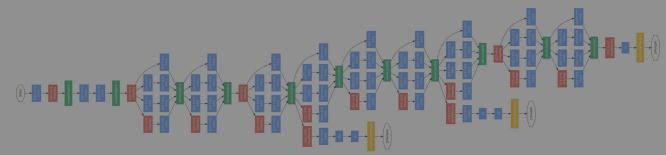
# TensorFlow handles complexity for you ...



Heterogenous  
systems



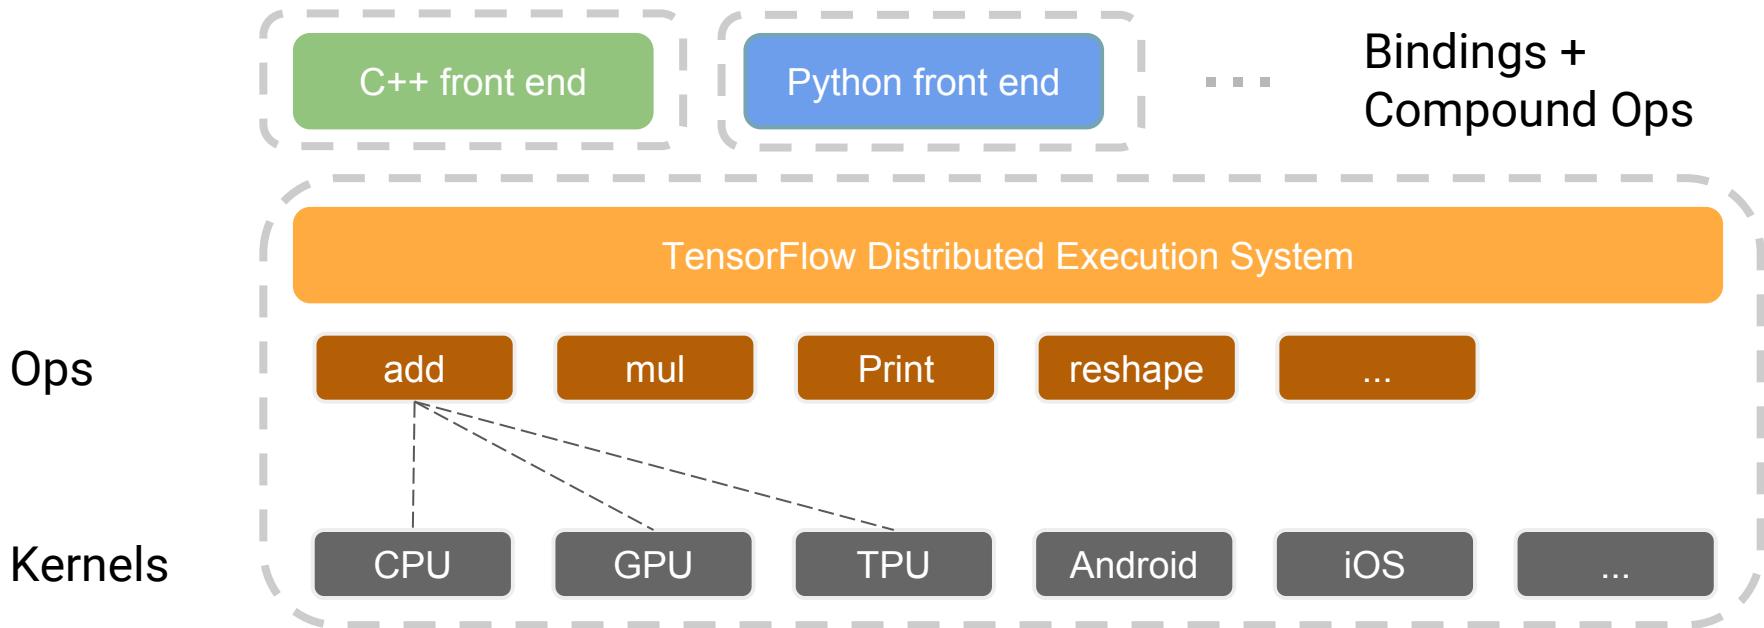
Distributed  
systems



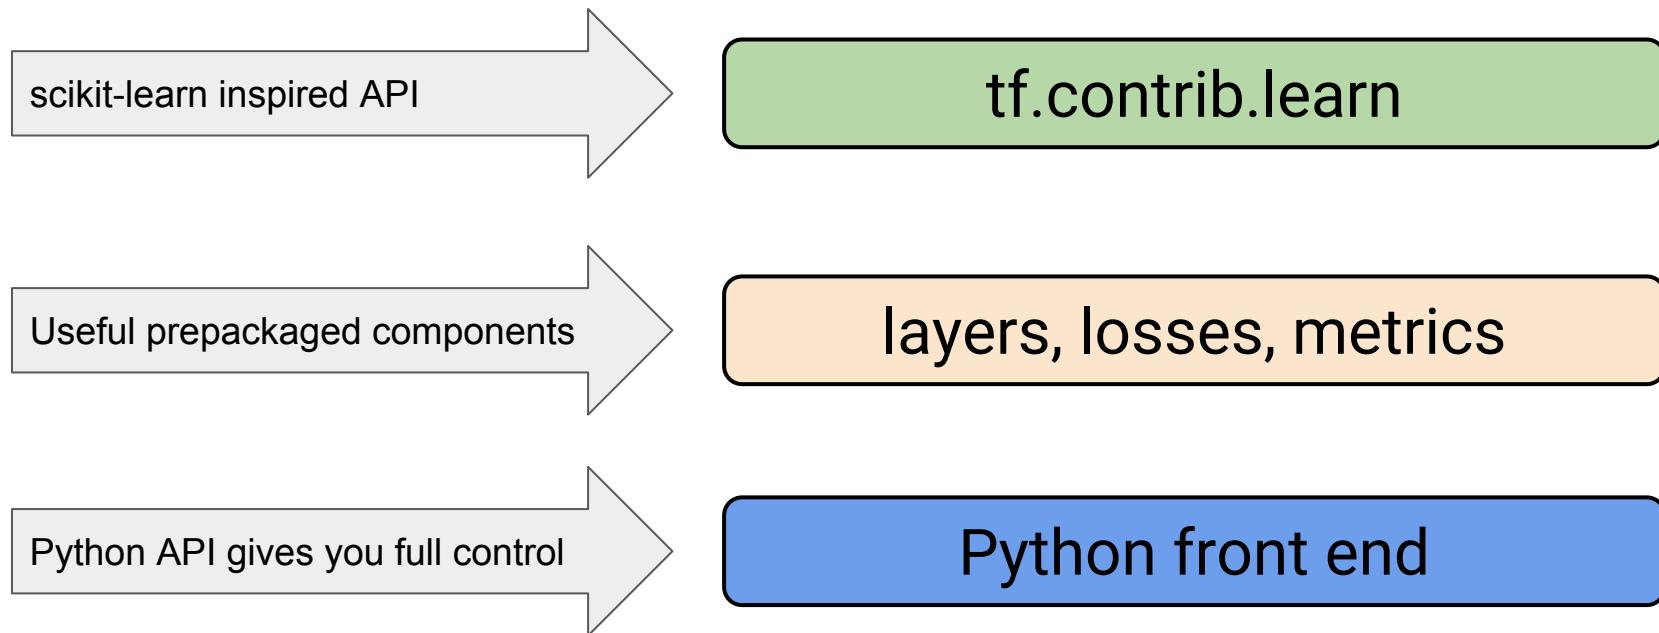
Complex models

... so you can focus on your ideas

# Architecture



# Toolkit hierarchy



# Community contributions

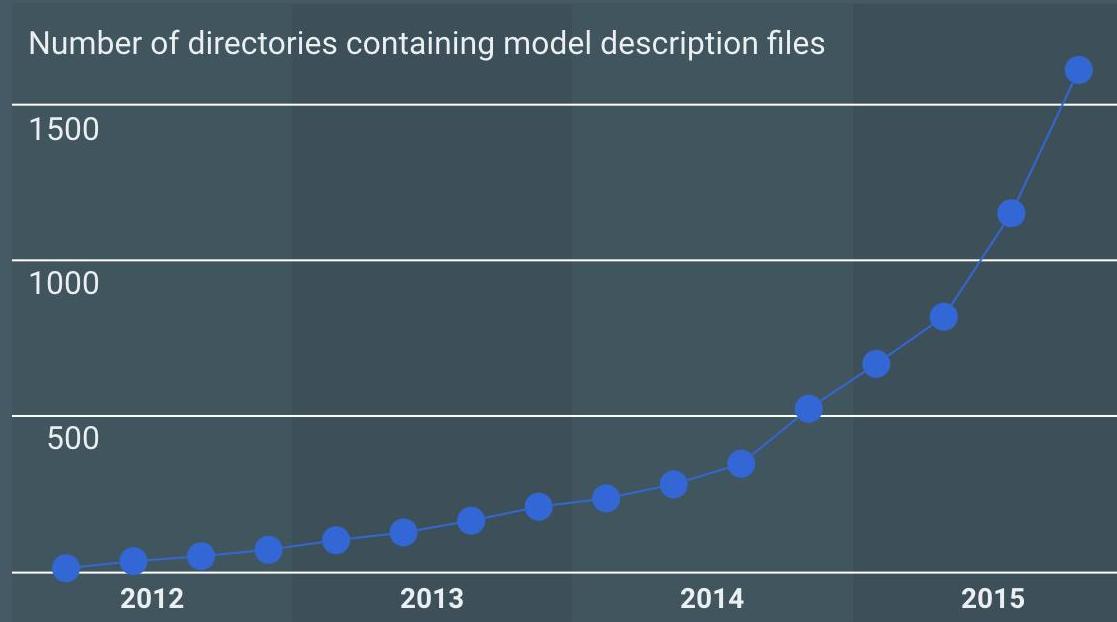


**Lab #1 Warm up**

**[goo.gl/nrdsxM](http://goo.gl/nrdsxM)**

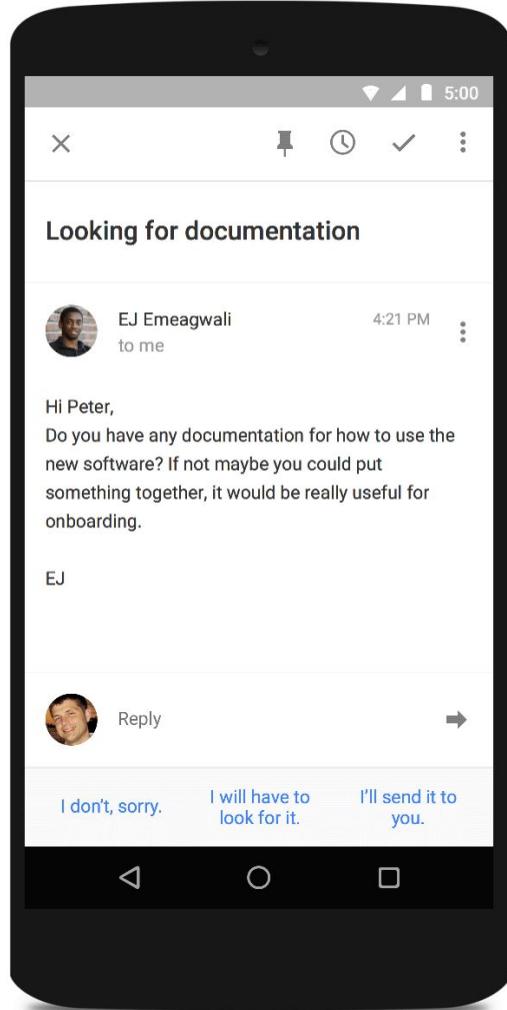
# TensorFlow at Google

# Growing use of deep learning at Google



Across many areas

- AlphaGo
- Apps
- Maps
- Photos
- Gmail
- Speech
- Android
- YouTube
- Translation
- Robotics Research
- Image Understanding
- Natural Language Understanding
- Drug Discovery



## Smart reply in Inbox by Gmail

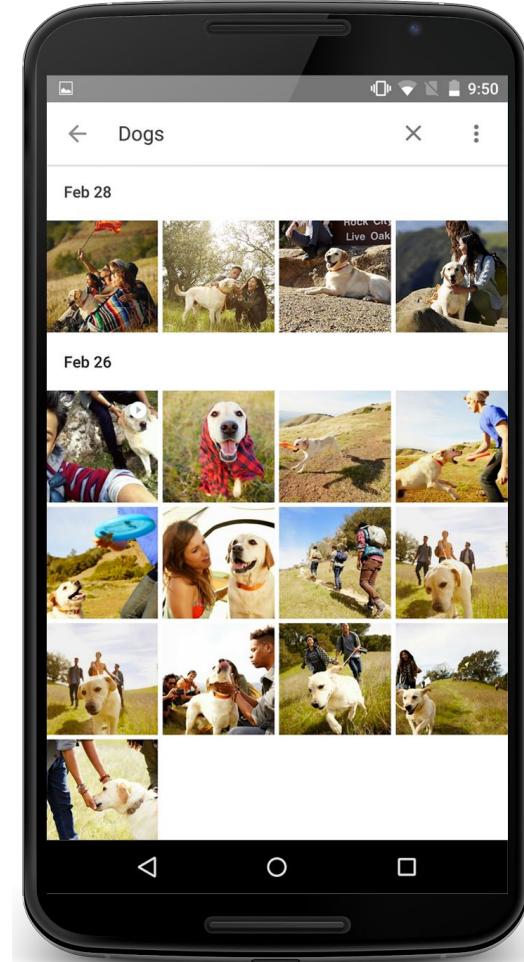
10%

of all responses  
sent on mobile



“Ok Google... Dog videos”







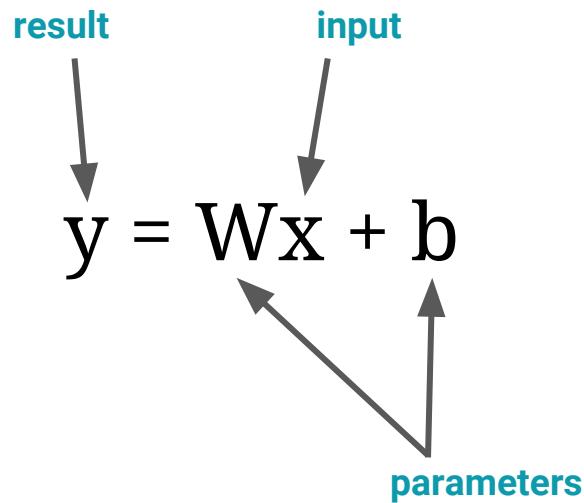
EXIT

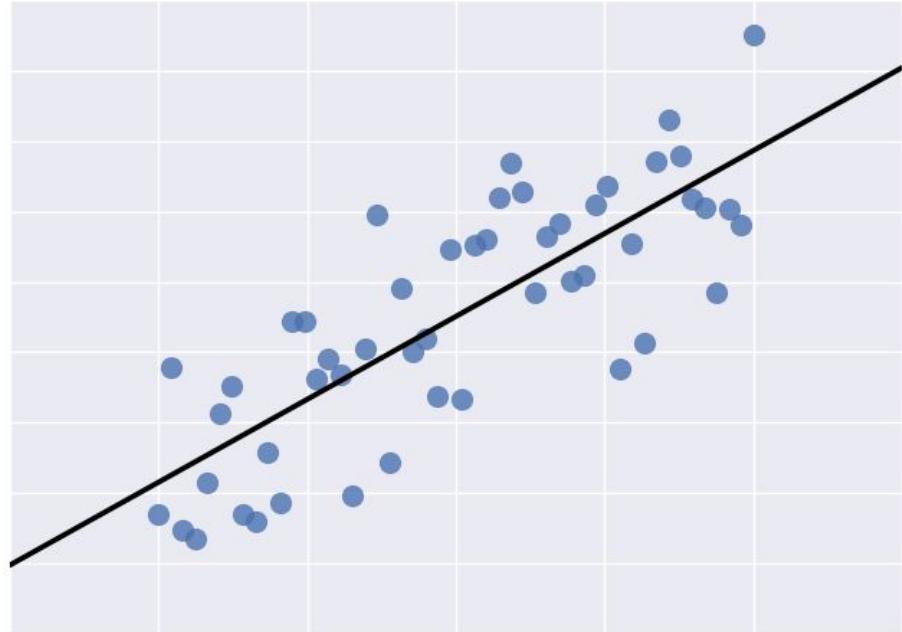
# Linear Regression

# Linear Regression

$$y = Wx + b$$

result      input  
parameters





# What are we trying to do?

**Mystery equation:**  $y = 0.1 * x + 0.3 + \text{noise}$

**Model:**  $y = W * x + b$

**Objective:** Given enough  $(x, y)$  value samples, figure out the value of  $W$  and  $b$ .

# Pattern

1. **Inference** (“code to predict y”)
2. **Loss** (“code to quantify how close our prediction was”)
3. **Eval** (“check your accuracy”)
4. **Train** (“code to update our variables to improve the next prediction”)

# Inference

```
import tensorflow as tf  
W = tf.get_variable(shape=[], name='W')  
b = tf.get_variable(shape=[], name='b')  
x = tf.placeholder(shape=[None],  
                   dtype=tf.float32,  
                   name='x')  
y = tf.matmul(W, x) + b  
  
with tf.Session() as sess:  
    sess.run(tf.initialize_all_variables())  
    print(sess.run(y, feed_dict={x: x_in}))
```

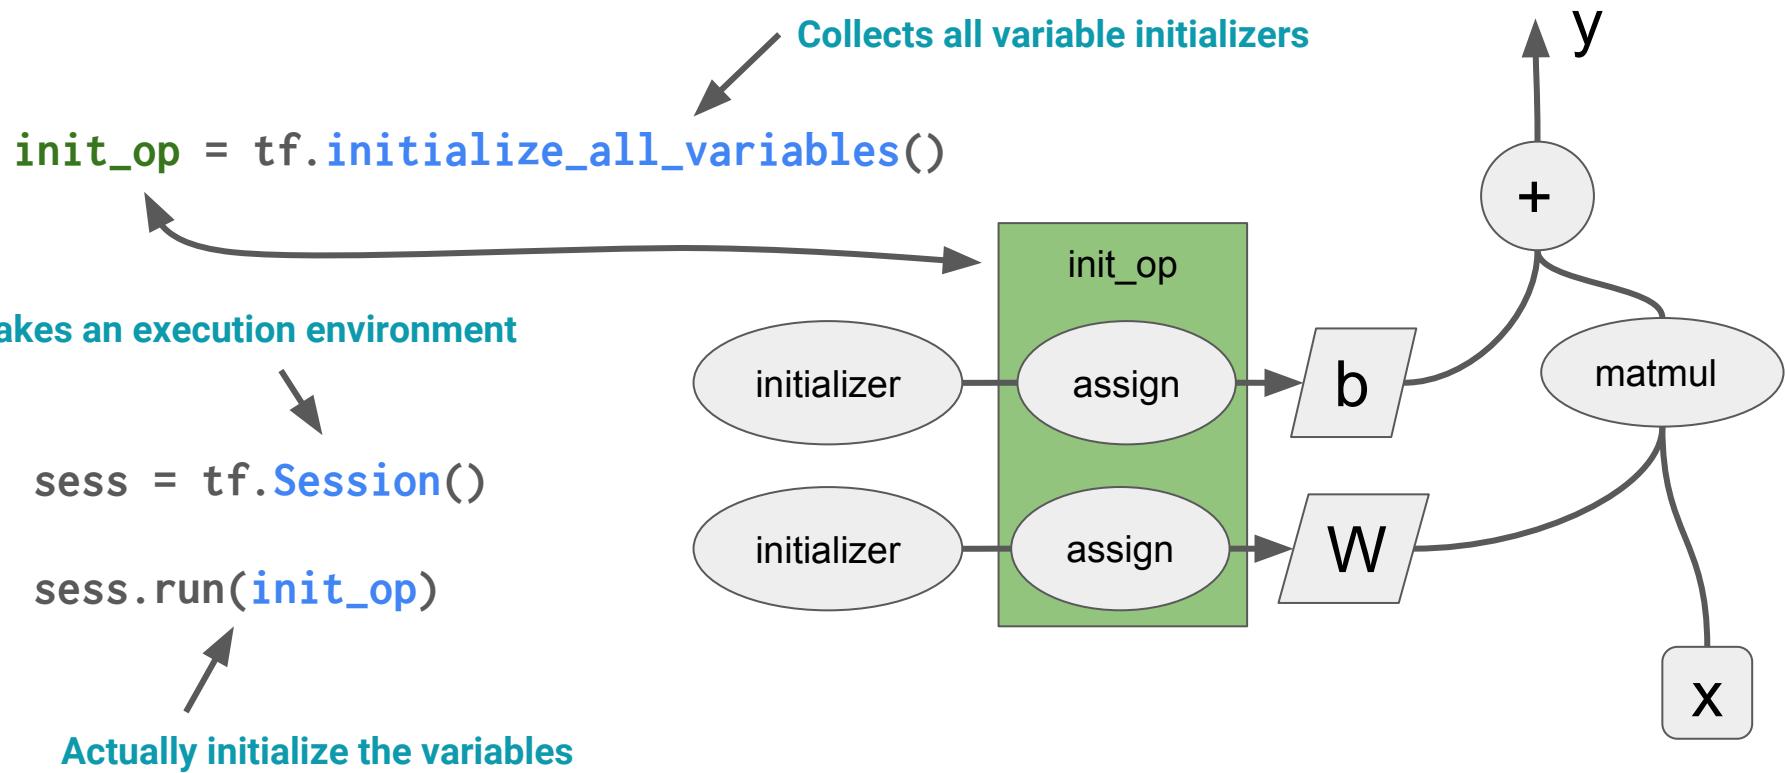
Build the graph

Prepare execution environment

Initialize variables

Run the computation (usually often)

# Variables Must be Initialized

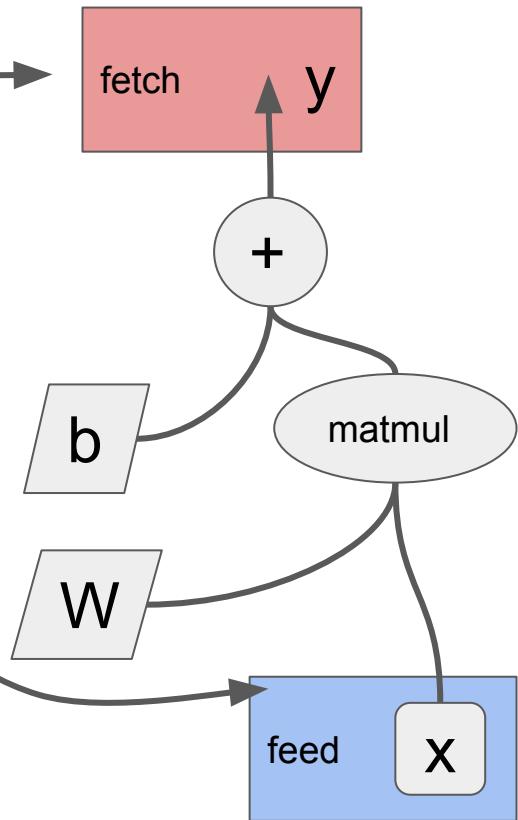


# Running the Computation

```
x_in = [3]
```

```
sess.run(y, feed_dict={x: x_in})
```

- Only what's used to compute a fetch will be evaluated
- All Tensors can be fed, but all placeholders must be fed



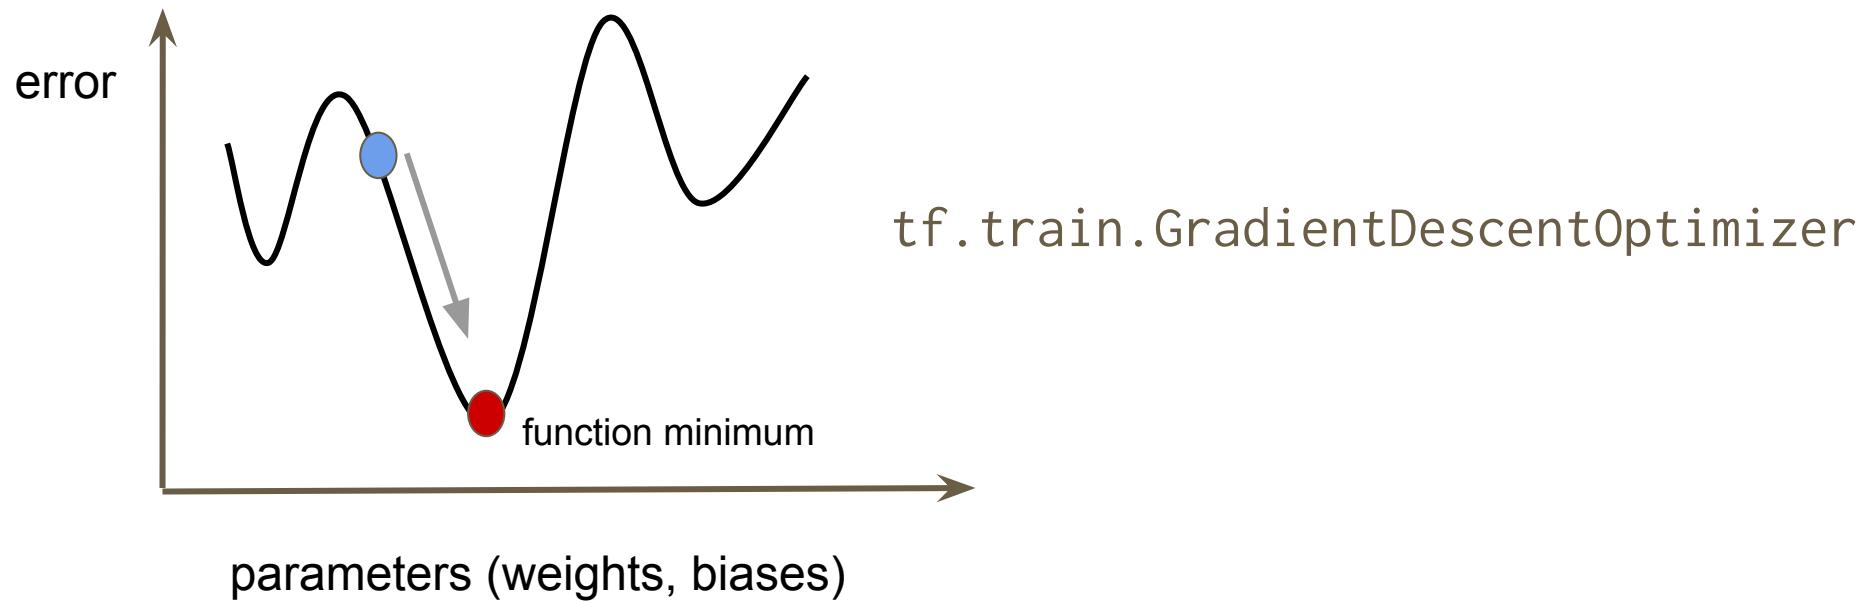
# Define a Loss

Given  $x$ ,  $y_{\text{label}}$ , compute a loss, for instance:

$$L = (y - y_{\text{label}})^2$$

```
# Create an operation that calculates loss.  
loss = tf.reduce_mean(tf.square(y - y_data))
```

# Minimize loss: optimizers



# Train

Feed  $(x, y_{\text{label}})$  pairs and adjust  $W$  and  $b$  to decrease the loss.

$$W \leftarrow W - \eta ( dL/dW )$$

$$b \leftarrow b - \eta ( dL/db )$$

TensorFlow computes  
gradients automatically

```
# Create an optimizer
```

```
optimizer = tf.train.GradientDescentOptimizer(0.5)
```

```
# Create an operation that minimizes loss.
```

```
train = optimizer.minimize(loss)
```

Learning rate

# Building graphs looks *mostly* like numpy

With special functions for deep learning

Numpy	TensorFlow
add	add
mul	mul
matmul	matmul
...	...
sum	reduce_sum
...	...
	sigmoid
	relu
	...

# 2\_linear\_regression\_model.ipynb

1. Import libraries (cell 1.1)
2. Create input data (cell 1.2)
3. Build inference graph (cell 1.3)
  - a. Create Variables to hold weights and biases.
  - b. Create Operations that produce logistic outputs.
4. Build training graph (cell 1.4)
  - a. Loss
  - b. Optimizer
  - c. Train\_op: Operation that minimizes Loss
5. Create session and run initialization (cell 1.5)
6. Perform training (cell 1.6)

# Exercises to run

1. Run all the cells
2. Uncomment 1.5 and see the graph
3. Uncomment 1.8 and compare values
4. Try to train a different linear function
  - a. If you can't visualize the values, try changing pylab.ylim()
5. Try a different initialization function.
  - a. tf.random\_normal
6. Try a different optimizer
  - a. tf.AdagradOptimizer

# Lab #2 Linear Regression

[goo.gl/nrdsxM](http://goo.gl/nrdsxM)

**tf.learn**

# That seemed complicated...

```
import tensorflow as tf

R = tf.contrib.learn.LinearRegressor(feature_columns=[  
    tf.contrib.layers.real_valued_column('', dimension=5)])  
  
R.fit(x=data[:,0:5], y=target, batch_size=100, max_steps=1000)  
  
R.evaluate(x=eval_data[:,0:5], y=target_eval  
  
R.predict(x=np.asarray([1.5, 3.4, 1.6, 1.9, 2.0]))
```

# tf.learn

- High level API for TensorFlow
- Implements scikit-learn API
- One-liners create the graph for us behind the scenes.

```
classifier = learn.DNNClassifier(hidden_units=[10,20,10],n_classes=3)
```

# Linear Regression with Taxi Data

## Problem

- Predict tip amount based on distance traveled (and other features)

## Solution

- Prepare a numpy ndarray.
- Each column represents a feature
- Last column is the target
- Use TF.Learn to train a Linear and Deep Regressor; compare the results





**Lab #2b TF.Learn**  
**[goo.gl/nrdsxM](http://goo.gl/nrdsxM)**

# Linear vs Nonlinear classifiers



# Tinker With a Neural Network Right Here in Your Browser.

## Don't Worry, You Can't Break It. We Promise.



Iterations  
000,000

Learning rate  
0.03

Activation  
Tanh

Regularization  
None

Regularization rate  
0

Problem type  
Classification

### DATA

Which dataset do you want to use?



Ratio of training to test data: 50%

Noise: 0

Batch size: 10

### FEATURES

Which properties do you want to feed in?

$X_1$



$X_2$



$X_1^2$



$X_2^2$



$X_1 X_2$



$\sin(X_1)$



2 HIDDEN LAYERS



2 neurons



The outputs are mixed with varying weights, shown by the thickness of the lines.

4 neurons

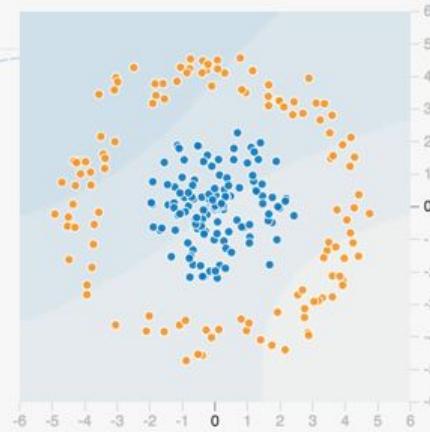


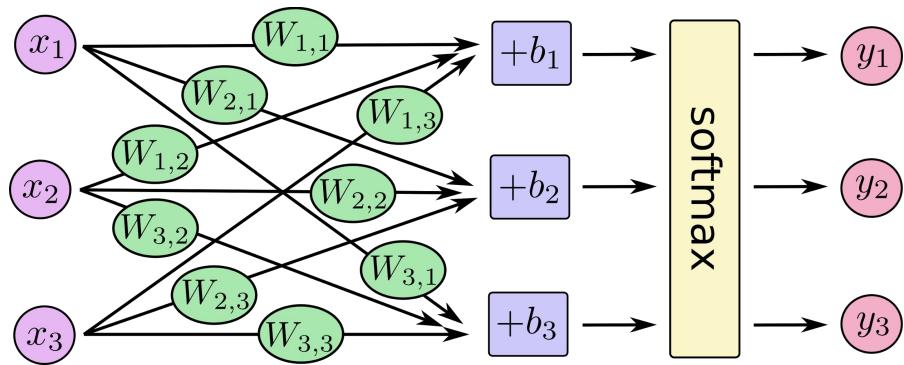
This is the output from one neuron.  
Hover to see it larger.

### OUTPUT

Test loss 0.513

Training loss 0.507



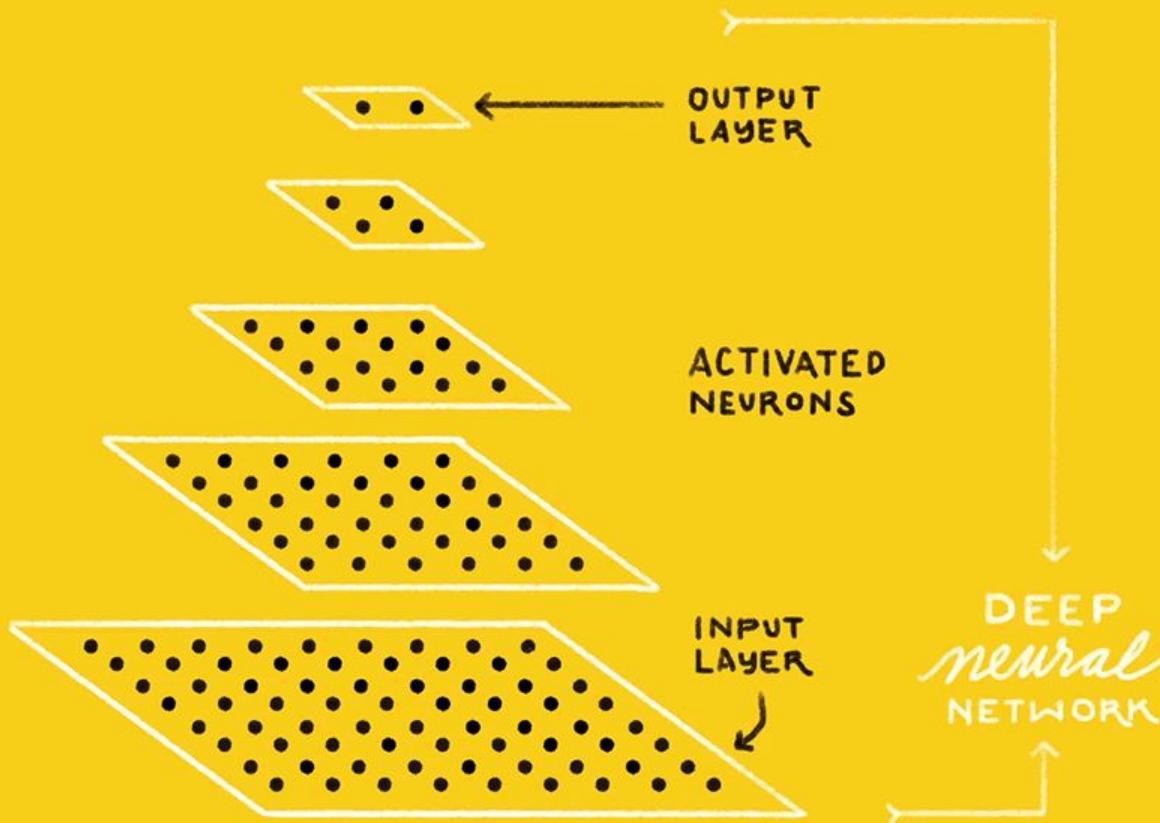


$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix} = \text{softmax} \begin{pmatrix} W_{1,1}x_1 + W_{1,2}x_2 + W_{1,3}x_3 + b_1 \\ W_{2,1}x_1 + W_{2,2}x_2 + W_{2,3}x_3 + b_2 \\ W_{3,1}x_1 + W_{3,2}x_2 + W_{3,3}x_3 + b_3 \end{pmatrix}$$

IS THIS A  
**CAT or DOG?**



CAT   DOG



# Stanford's cs231n

# Convolutional Neural Networks for Visual Recognition

## CS231n Convolutional Neural Networks for Visual Recognition

These notes accompany the Stanford CS class [CS231n: Convolutional Neural Networks for Visual Recognition](#). For questions/concerns/bug reports regarding contact [Justin Johnson](#) regarding the assignments, or contact [Andrej Karpathy](#) regarding the course notes. You can also submit a pull request directly to our [git repo](#). We encourage the use of the [hypothes.is](#) extension to annotate comments and discuss these notes inline.

### Winter 2016 Assignments

[Assignment #1: Image Classification, kNN, SVM, Softmax, Neural Network](#)

[Assignment #2: Fully-Connected Nets, Batch Normalization, Dropout, Convolutional Nets](#)

[Assignment #3: Recurrent Neural Networks, Image Captioning, Image Gradients, DeepDream](#)

### Module 0: Preparation

[Python / Numpy Tutorial](#)

[IPython Notebook Tutorial](#)

[Terminal.com Tutorial](#)

[AWS Tutorial](#)

### Module 1: Neural Networks

[Image Classification: Data-driven Approach, k-Nearest Neighbor, train/val/test splits](#)

[L1/L2 distances, hyperparameter search, cross-validation](#)

[Linear classification: Support Vector Machine, Softmax](#)

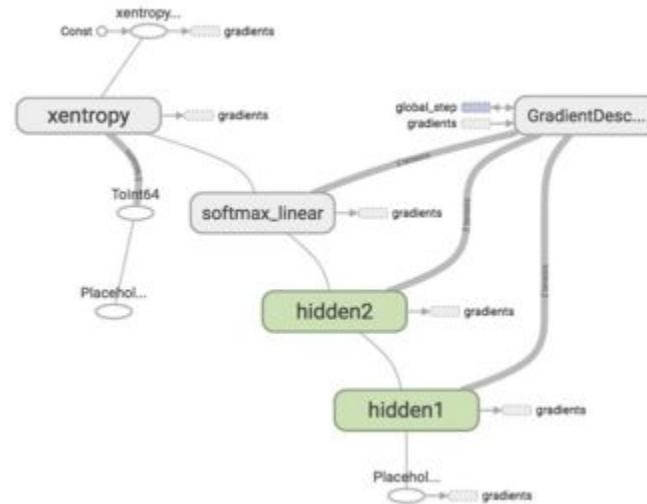
# Lab #3 Comparing Classifiers

[goo.gl/nrdsxM](http://goo.gl/nrdsxM)

# MNIST

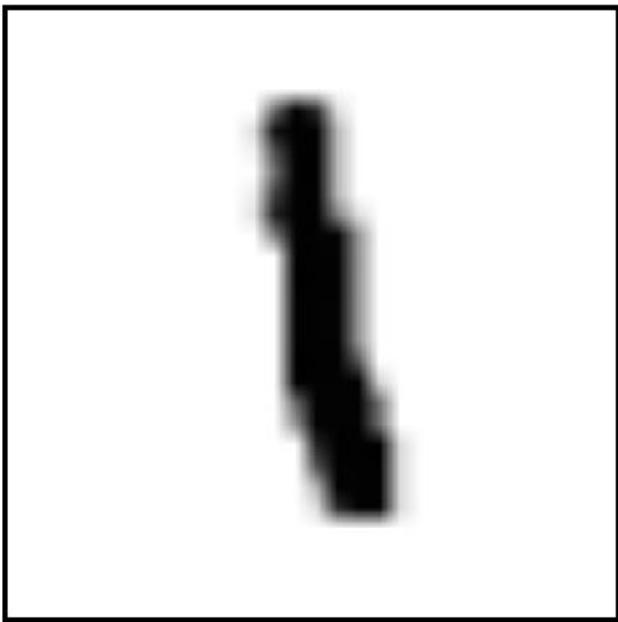


# What are we trying to do?



**Objective:** Given enough images and labels, figure out the weights and biases so the model can correctly identify digits.

# What we see



## What the computer “sees”

2

# 3\_mnist.ipynb

1. Import libraries (cell 2.1)
2. Define a set of constants for convenience (cell 2.2)
3. Get input data (cell 2.3)
4. Build inference graph (cell 2.4)
5. Build training graph (cell 2.5)
6. Build the complete graph for training, and saving checkpoints (cell 2.6)
7. Run training for MAX\_STEPS and save checkpoint at the end (cell 2.7)
8. Run evaluation based on the saved checkpoint (cell 2.8)

# Exercises to run

1. Run all the cells
2. Uncomment the graph writing to see what the graph looks like.
3. Print out loss value every 100 steps. What do you notice?
4. Save checkpoints every 500 steps.
5. Run evaluation with different checkpoints. What do you notice?
6. Run evaluation on the complete validation set.
7. Build evaluation graph from scratch instead of importing from meta graph.
8. Experiment with different optimizers

# Lab #4 Deep MNIST

## [goo.gl/nrdsxM](http://goo.gl/nrdsxM)

# Convolutional Neural Networks

1 <small><math>\times 1</math></small>	1 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	0	0
0 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	1 <small><math>\times 0</math></small>	1	0
0 <small><math>\times 1</math></small>	0 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved Feature

Convolution with 3x3 Filter. Source:

[http://deeplearning.stanford.edu/wiki/index.php/Feature\\_extraction\\_using\\_convolution](http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution)

### Single depth slice

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

x

y

max pool with 2x2 filters  
and stride 2

6	8
3	4

Max pooling in CNN. Source: <http://cs231n.github.io/convolutional-networks/#pool>, via  
<http://www.wildml.com/2015/11/understanding-convolutional-neural-networks-for-nlp/>

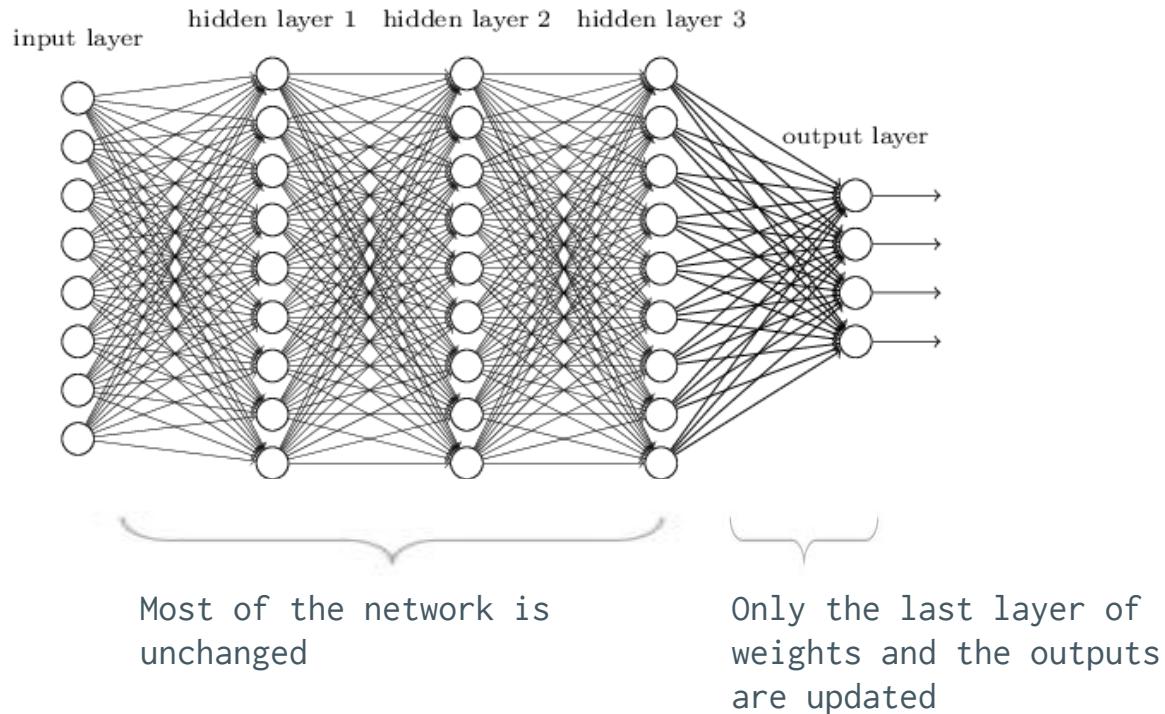
[bit.ly/tensorflow-workshop](http://bit.ly/tensorflow-workshop)

# Lab #5

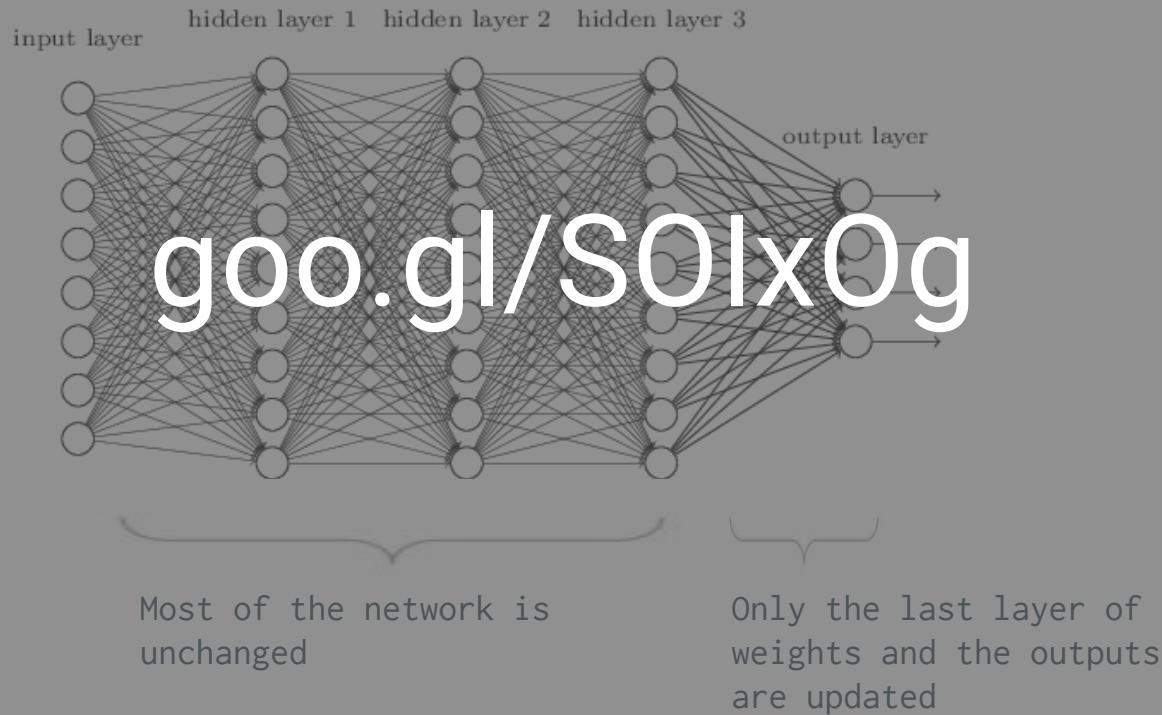
## Convolutional MNIST at TensorFlow.org

# Image Retraining

# Image Retraining



# Image Retraining



# TensorFlow for Poets

The screenshot shows the TensorFlow For Poets codelab interface. On the left, a vertical sidebar lists steps from 1 to 9: Introduction, Setting Up, Installing and Running the TensorFlow Docker Image, Retrieving the images, (Re)training Inception, Using the Retrained Model, Optional Step: Trying Other Hyperparameters, Optional Step: Training on Your Own Categories, and Next Steps. Step 1 is highlighted. At the bottom of the sidebar, it says "Did you find a mistake? Please [file a bug](#)". The main content area is titled "1. Introduction". It describes TensorFlow as an open source library for numerical computation, specializing in machine learning applications. It explains that in this lab, users will learn how to install and run TensorFlow on a single machine, and will train a simple classifier to classify images of flowers. It mentions that transfer learning will be used, starting with an Inception v3 model trained on ImageNet. It also lists what users will learn (installing TensorFlow, using Bazel and Python to train an image classifier, classifying images with a trained classifier) and what they need (basic Unix commands, a fast computer, and time). A note at the bottom states that the codelab does not cover Windows, but links to success stories for other platforms.



Codelab  
[goo.gl/xGsB9d](https://codelabs.developers.google.com/codelabs/tensorflow-for-poets/index.html)

Video  
[Machine Learning Recipes #6](https://www.youtube.com/watch?v=KJLjyfzXWUo)

# TensorFlow for Poets

The screenshot shows the TensorFlow For Poets codelab interface. On the left, a vertical navigation bar lists 9 steps: 1. Introduction (selected), 2. Setting Up, 3. Installing and Running the TensorFlow Docker Image, 4. Retrieving the images, 5. (Re)training Inception, 6. Using the Retrained Model, 7. Optional Step: Trying Other Hyperparameters, 8. Optional Step: Training on Your Own Categories, and 9. Next Steps. The main content area is titled "1. Introduction". It starts with a paragraph about TensorFlow being an open source library for numerical computation, specializing in machine learning applications. It explains that in this codelab, you will learn how to install and run TensorFlow on a single machine, and will train a simple classifier to classify images of flowers. Below this, there's a section titled "What are we going to be building?" which states that in this lab, they will be using transfer learning, starting with a model trained on the ImageNet Large Visual Recognition Challenge. It then describes the optional steps for trying other hyperparameters and training on your own categories. A "What you will learn" section lists: How to install and run TensorFlow Docker images, How to use Bazel and Python to train a image classifier, and How to classify images with your trained classifier. A "What you need" section lists: A basic understanding of Unix commands, A fast computer running OS X or Linux, and A fair amount of time. At the bottom, it notes that the codelab does not cover Windows and provides a link for those who have had success getting the Docker image working.



Codelab

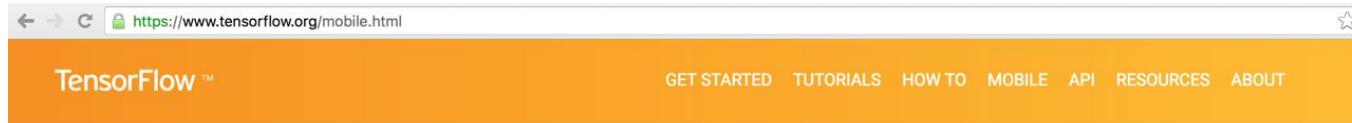
[goo.gl/xGsB9d](https://codelabs.developers.google.com/codelabs/tensorflow-for-poets/index.html)

Video

[Machine Learning Recipes #6](#)

# Mobile TensorFlow

# Mobile TensorFlow



A screenshot of a web browser displaying the TensorFlow mobile website at <https://www.tensorflow.org/mobile.html>. The page has a yellow header with the TensorFlow logo and navigation links for GET STARTED, TUTORIALS, HOW TO, MOBILE, API, RESOURCES, and ABOUT. The main content area features a section titled "Mobile TensorFlow" with text about its design for mobile platforms and sample code for Android, iOS, and Raspberry Pi. To the right is an image of a smartphone showing a cat being identified with a confidence score of 15% Tabby.

## Mobile TensorFlow

TensorFlow was designed with mobile and embedded platforms in mind.

We have sample code and build support you can try now for these platforms:

[Android](#)

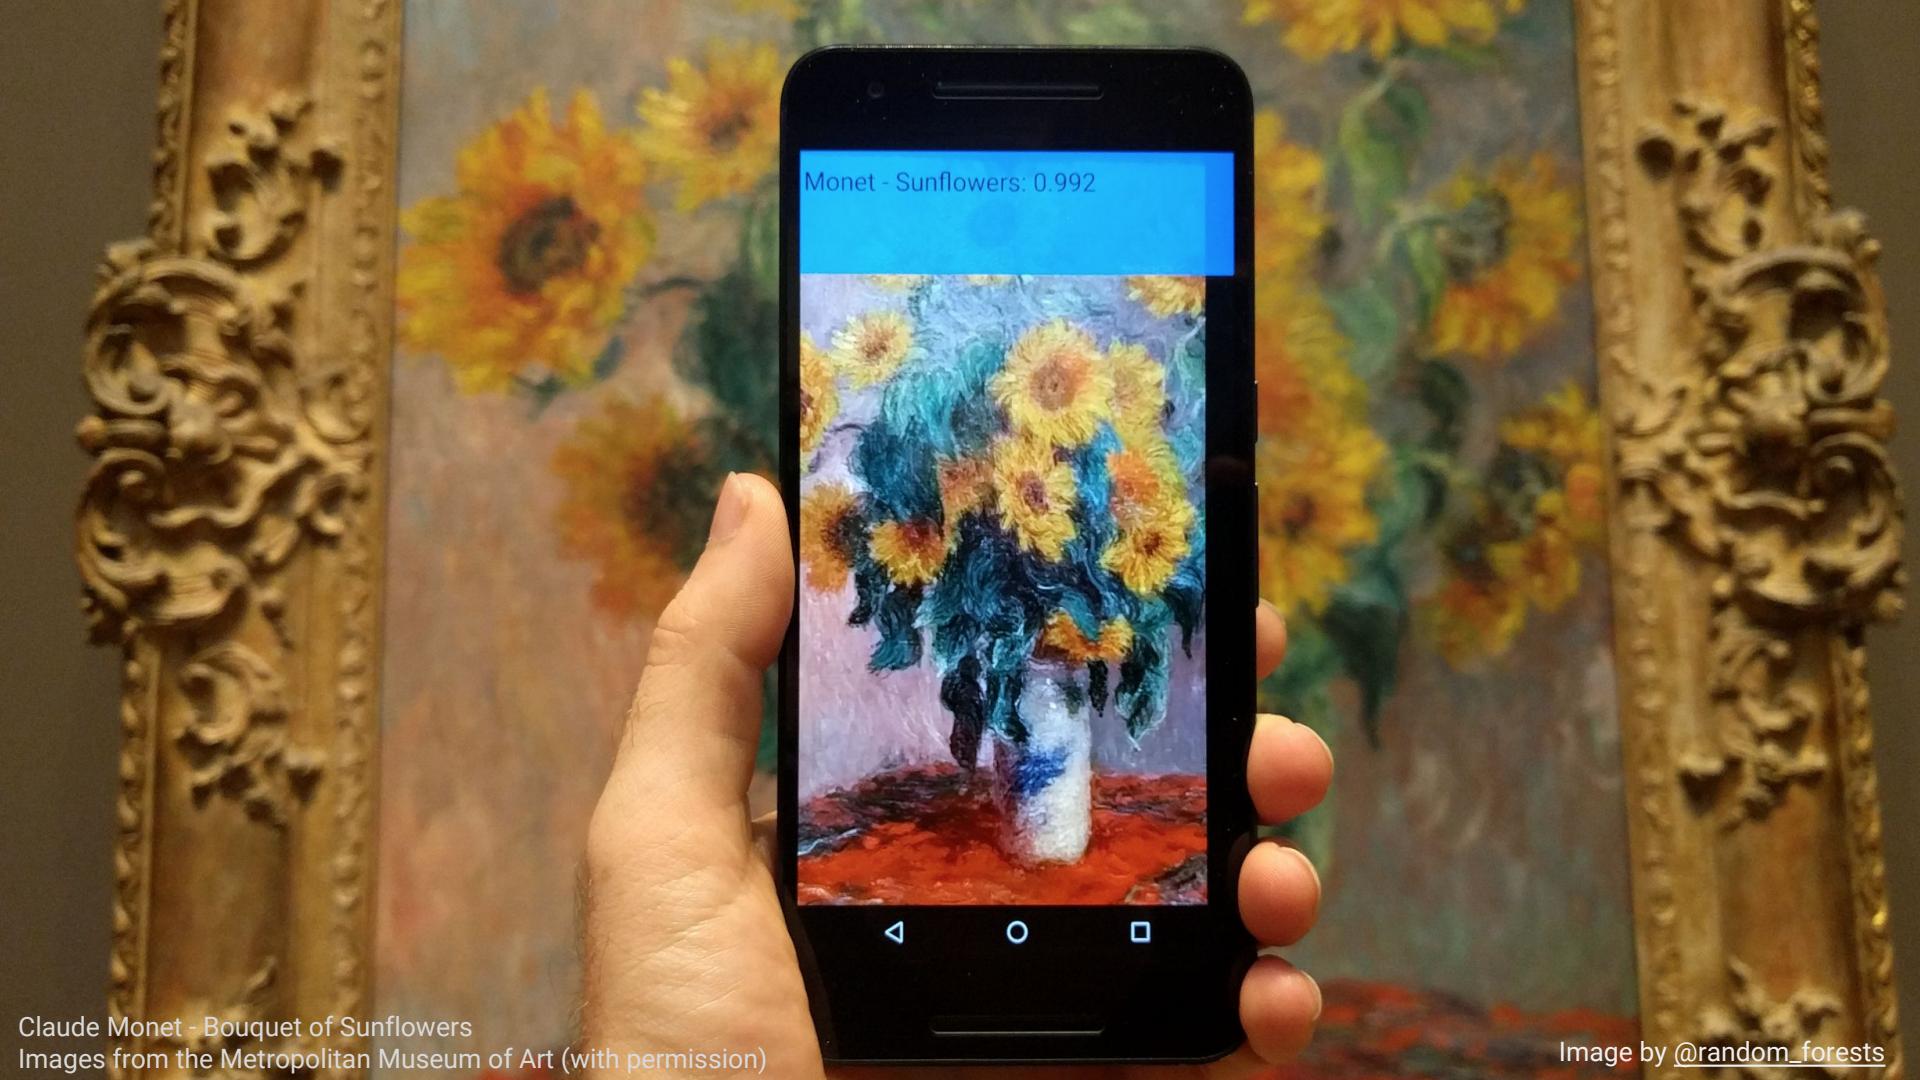
[iOS](#)

[Raspberry Pi](#)

Many applications can benefit from on-device processing. Google Translate's instant visual translation is a great example. By running its processing locally, users get an incredibly responsive and interactive experience.

Mobile TensorFlow makes sense when there is a poor or missing network connection, or where sending continuous data to a server would be too expensive. We are working to help developers make lean mobile apps using TensorFlow, both by continuing to reduce the code footprint, and supporting [quantization](#) and [lower precision arithmetic](#) that reduce model size.





Claude Monet - Bouquet of Sunflowers  
Images from the Metropolitan Museum of Art (with permission)

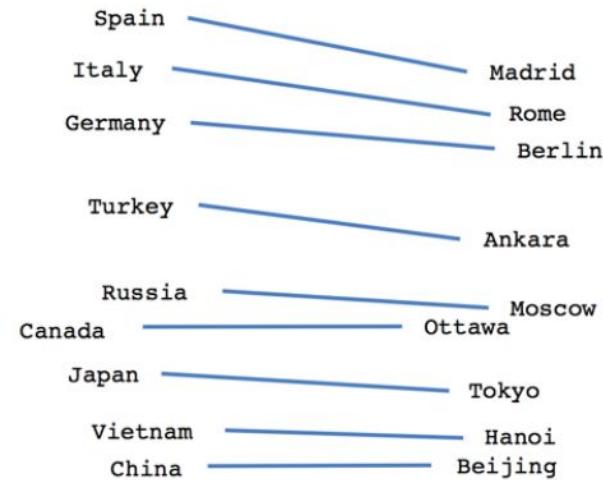
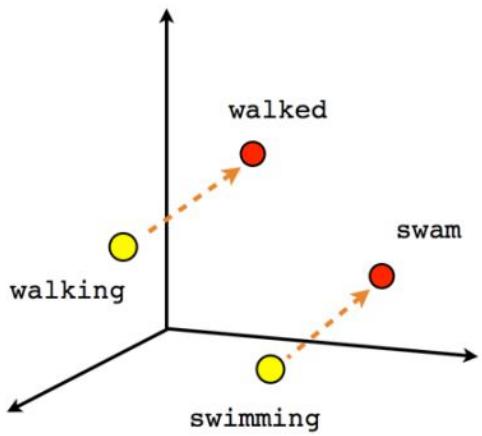
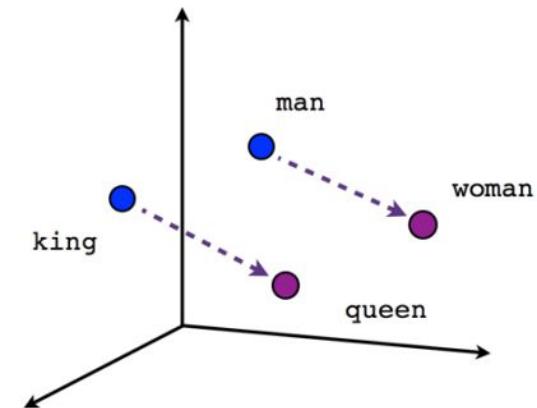
Image by [@random\\_forests](#)

# Lab #6 TensorFlow for Poets

# Word Embeddings

Nearest to government:

'governments', 'leadership',  
'regime', 'crown', 'rule',  
'leaders', 'parliament',  
'elections', ...



# Lab #7 Word embeddings

[goo.gl/nrdsxM](http://goo.gl/nrdsxM)

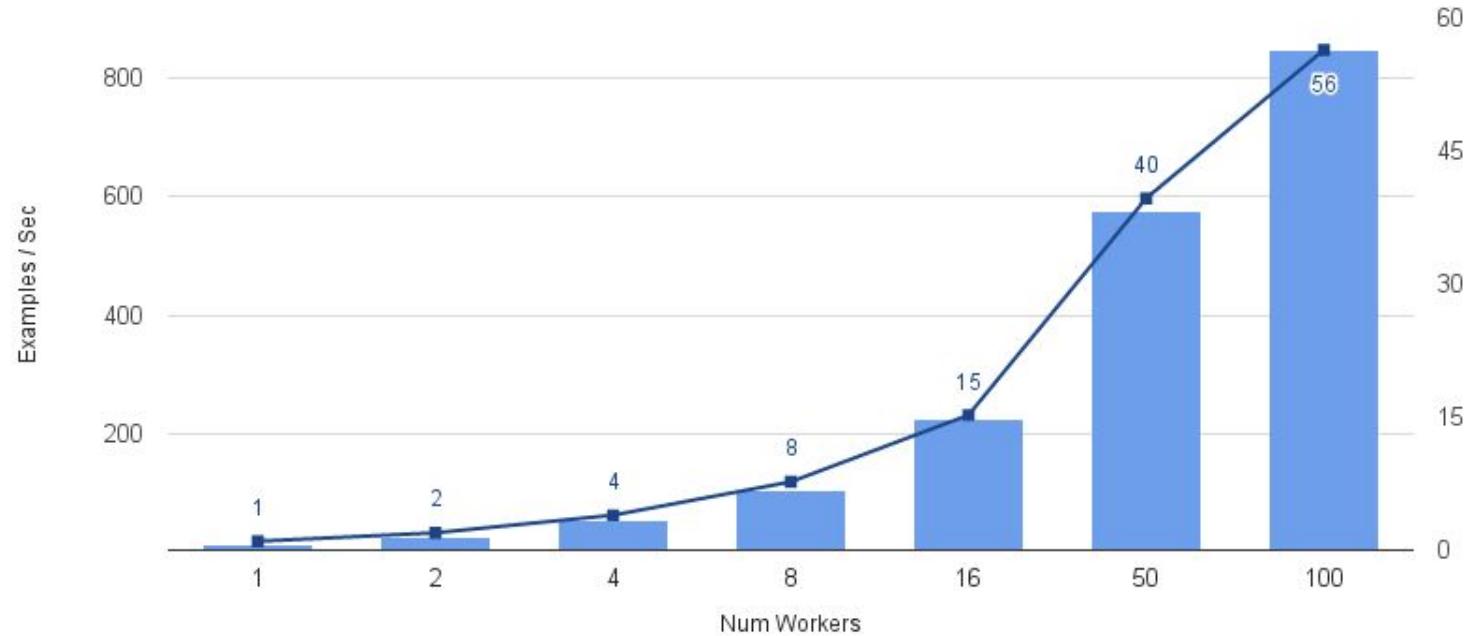
Also check [tensorflow.org](http://tensorflow.org)  
for analogies

# Lab #8 Document classification

[https://github.com/tensorflow/tensorflow/blob/master/tensorflow/examples/skflow/text\\_classification.py](https://github.com/tensorflow/tensorflow/blob/master/tensorflow/examples/skflow/text_classification.py)

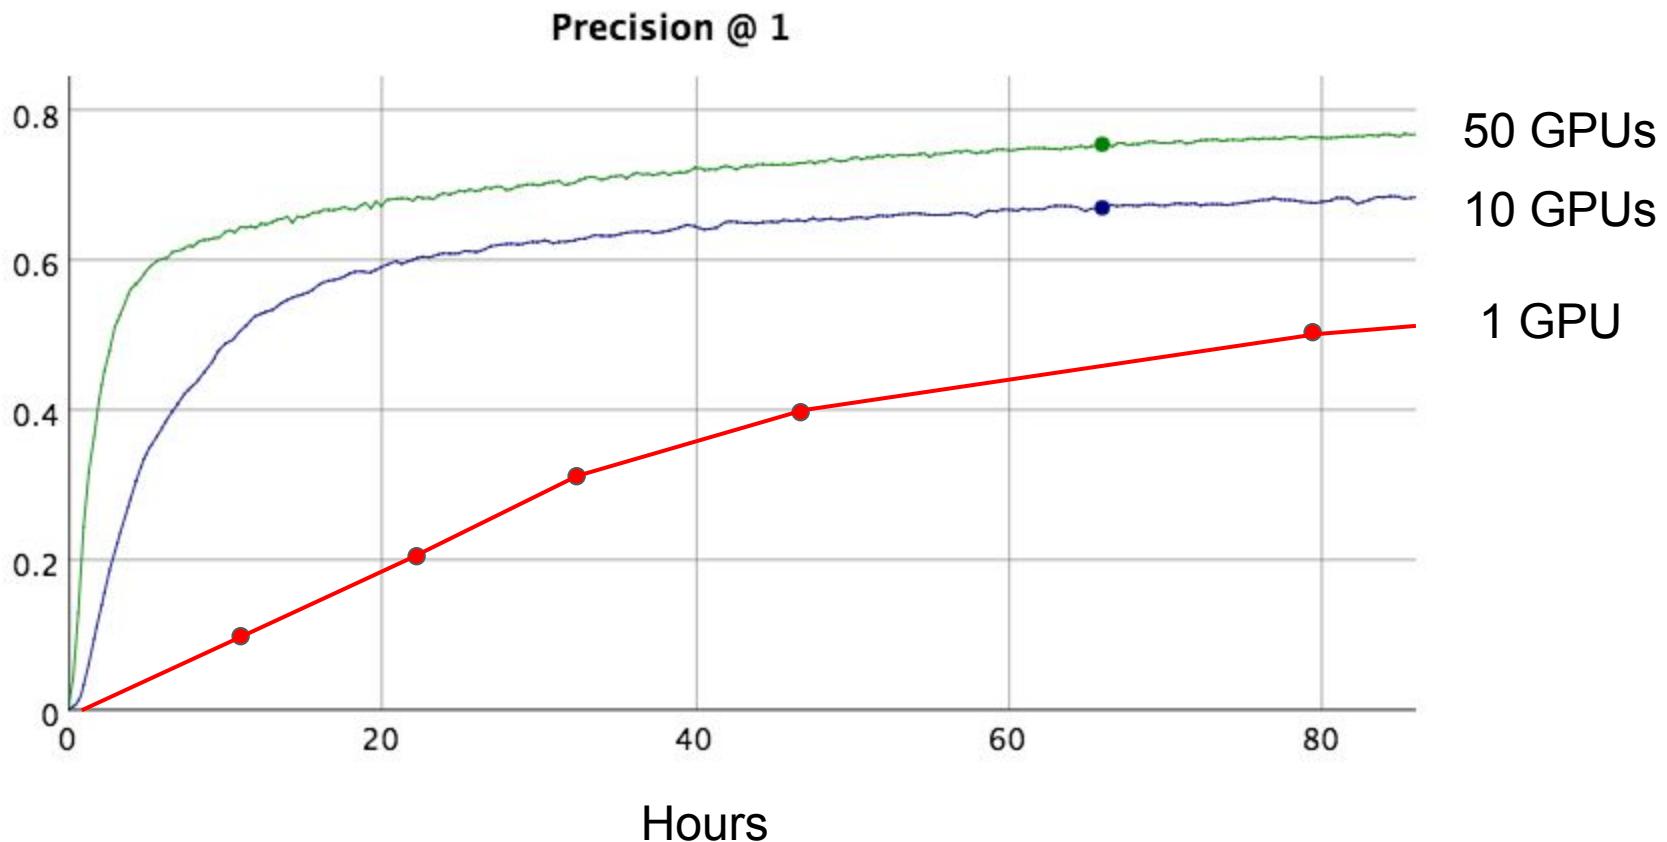
# Distributed TensorFlow

# TensorFlow

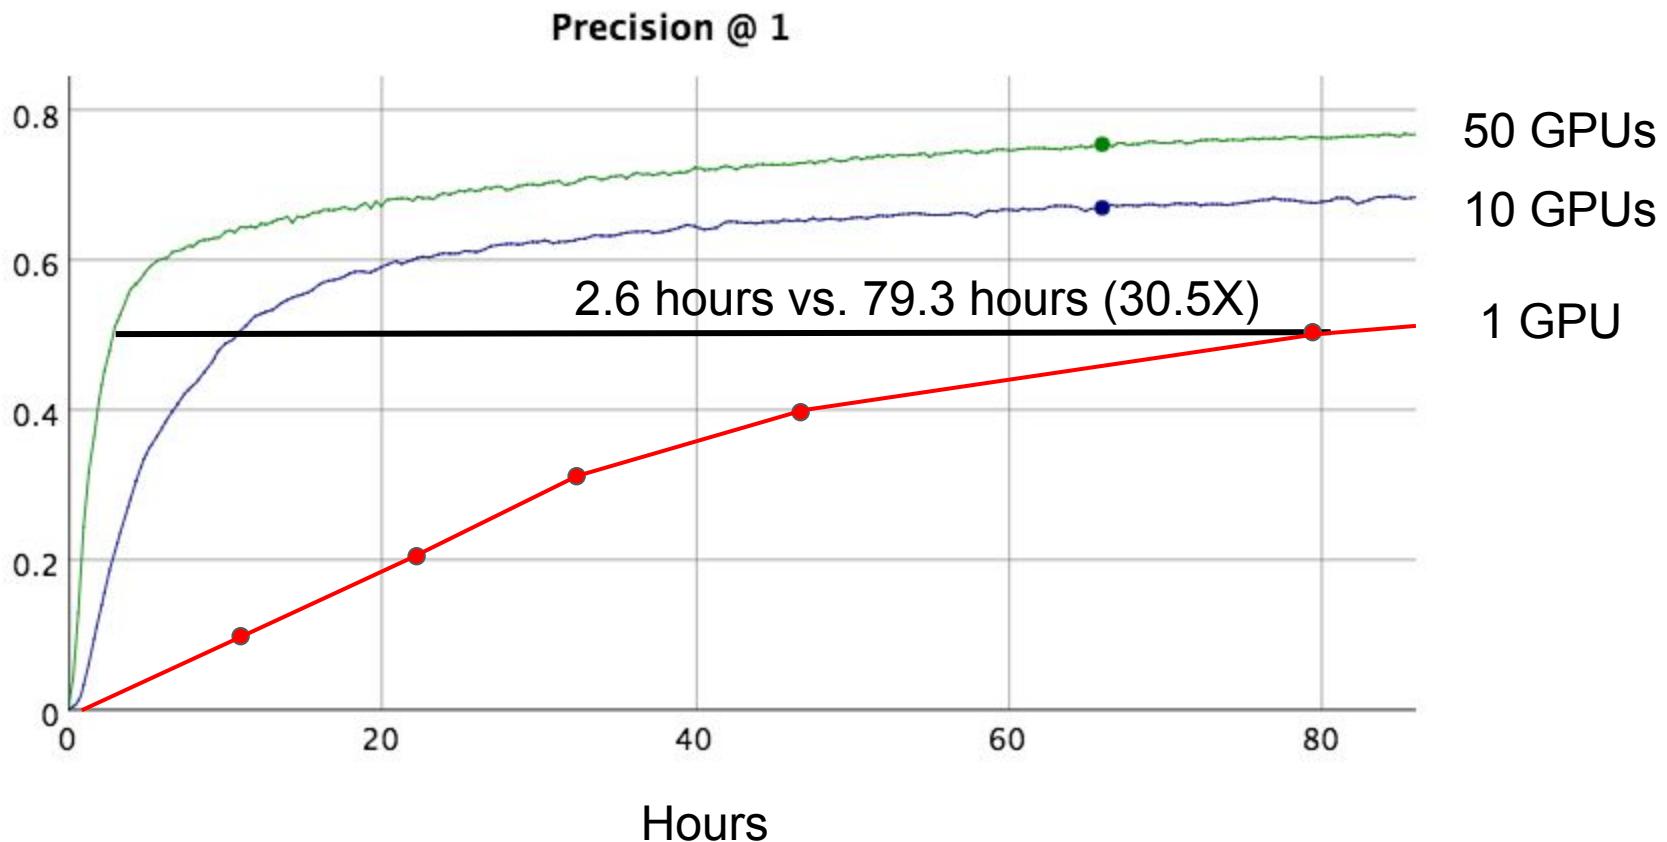


<https://arxiv.org/abs/1604.00981>

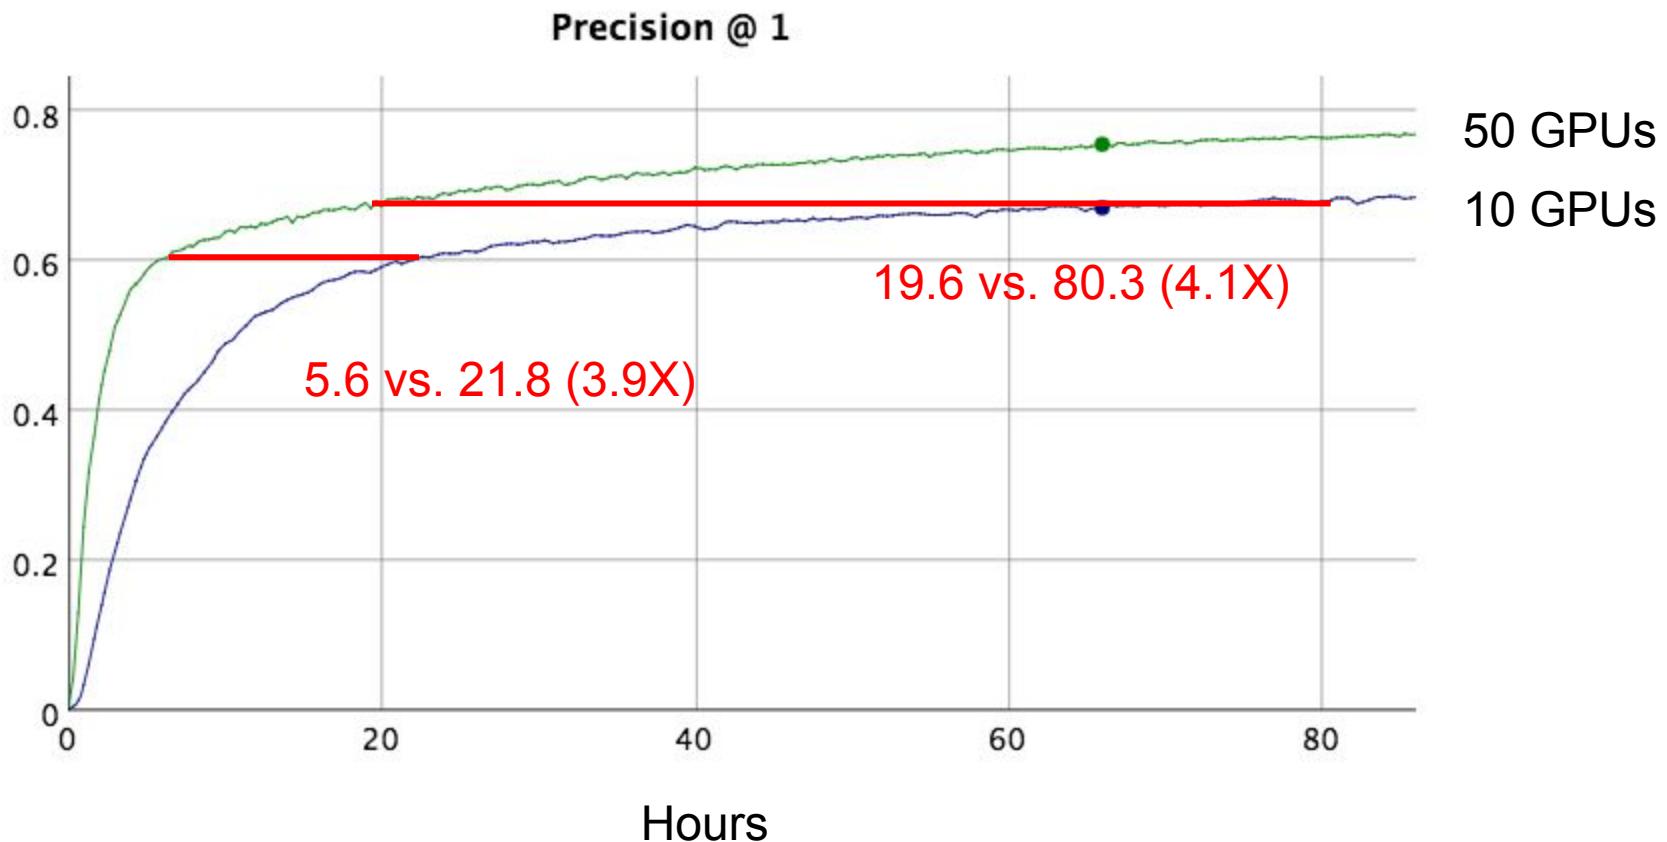
# Image Model (Inception) Synchronous Training



# Image Model (Inception) Synchronous Training



# Image Model (Inception) Synchronous Training



# Google Cloud Machine Learning Platform LIMITED PREVIEW

- Training and prediction at Scale
- Fully Managed Platform
- Powered by TensorFlow
- Seamless integration with Google Cloud Platform





**Bottomline: TensorFlow is  
for Machine Learning -- from  
lab to production.**

# Fun experiments



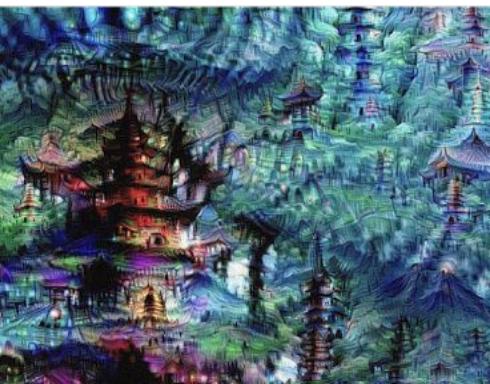
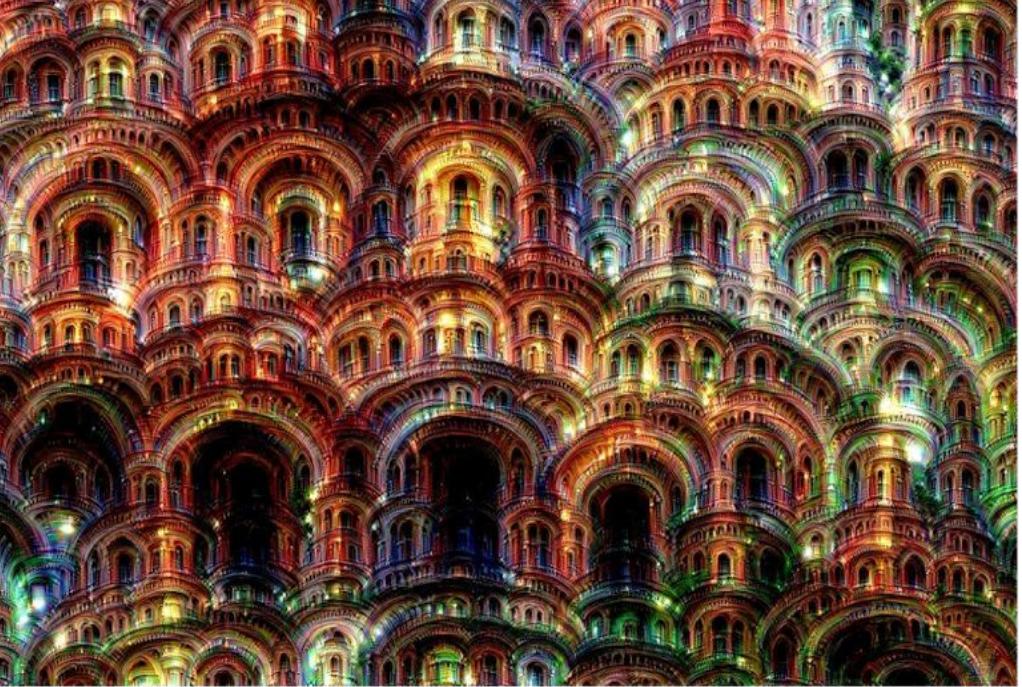
Image: Leonid Afremov – Rain Princess



Image: Stephen Wilkes, National Geographic

The background of the image is a colorful, abstract painting of a city. It features a dense cluster of skyscrapers in various colors like blue, green, yellow, and orange against a backdrop of a cloudy sky. In the foreground, there's a wide street filled with numerous yellow taxis and some people walking. The overall style is painterly and somewhat blurry.

goo.gl/fyDxhC





goo.gl/OCYseb





goo.gl/67hfSo

# Fun experiments

**Deep Dream: [goo.gl/OCYseb](http://goo.gl/OCYseb)**

**Style transfer: [goo.gl/fyDxhC](http://goo.gl/fyDxhC)**

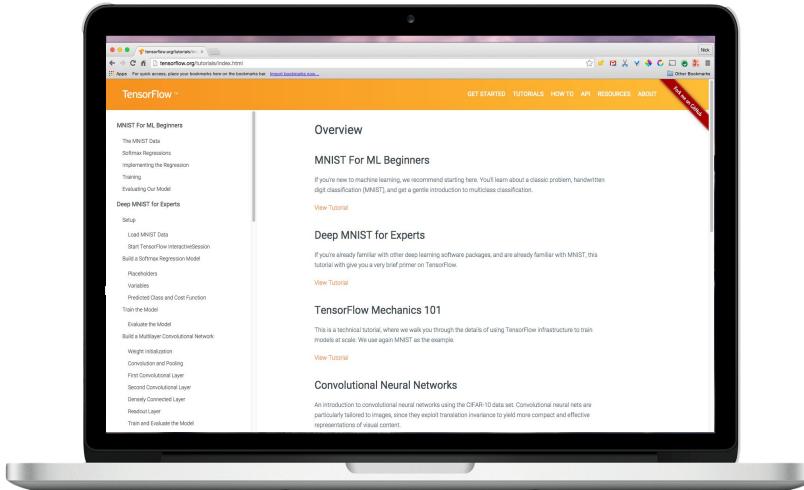
**Image completion: [goo.gl/67hfSo](http://goo.gl/67hfSo)**

# Next steps

Tutorials and code  
[tensorflow.org](http://tensorflow.org)

Want to learn more?  
Udacity class [goo.gl/iHsslI](http://goo.gl/iHsslI)

Totally new to ML?  
Recipes [goo.gl/KewA03](http://goo.gl/KewA03)



# Thanks and have fun!



Martin Wicke  
@martin\_wicke



Josh Gordon  
@random\_forests